

7. SQL

(Structured Query Language)

SQL

- Initially, SEQUEL language (System R) project at IBM
- Three relational languages : QUEL, QBE, SQL
- (Later renamed) Structured Query Language (SQL)
- Currently, Standard Relational database Language
 - Users need to be less concerned to migrate to other types of relational DBMS
- Very-High-level language
 - **Declarative** (= Non-procedural) language
 - Users do not need to specify how (= what order of) to execute query operations.
 - Say "what to do" rather than "how to do"
 - Relational algebra is a procedural language.

SQL

- Table (= relation), row (= tuple), column (= attribute)
- A table consists of :
 - Base Table
 - View (= Virtual table)
- Basic commands;
 - Data Definition Language (DDL)
 - Data Manipulation Language (DML)
 - Data Control Language (DCL)
- DDL is used for defining schemas, tables, and views.
 - CREATE, DROP, ALTER
- DML is used for retrieving and modifying tuples in a table.
 - SELECT (FROM WHERE)
 - INSERT, DELETE, UPDATE

CREATE TABLE

- A table is defined using **CREATE TABLE** command:

```
CREATE TABLE R  
(  
    A1 : D1  
    A2 : D2  
    . . . . .  
    An : Dn)  
    Constraints )
```

- R : relation name
 - A_i : attribute name
 - D_i : domain (=data type) of A_i
 - Constraints : Integrity Constraints
- A relation defined by CREATE TABLE is called a **"base table"**.
 - The relation (and its tuples) are "physically stored".
 - All created tables are initially "empty".
 - Attributes are ordered as they are specified in CREATE TABLE


CREATE TABLE : Attribute Data Types

- System Defined
 - Numeric : **INT, SMALLINT, FLOAT(or REAL), . . .**
 - Character String : **CHAR(n), VARCHAR(n), CLOB, . .**
 - Bit String : **BIT(n), BIT VARYING(n), BLOB, . .**
 - Boolean : **TRUE, FALSE, UNKNOWN**
 - DATE : **YEAR, MONTH, DAY, . . .**
 -
- User Defined
 - It is possible for users to specify domain of attribute directly;
 - For example, we can create domain "SSN-TYPE";
CREATE DOMAIN SSN-TYPE AS CHAR(9)
 - We can use this SSN-TYPE for defining attributes SSN, Super_SSN, Mger-SSN, . . . ;
 - What if we change later SSN data type to another one?

CREATE TABLE : Constraints

- Key Integrity
 - Primary key and keys can be specified by **PRIMARY KEY** and **UNIQUE**, respectively
- Entity Integrity
 - This must be specified by **NOT NULL** on **PRIMARY KEY**
 - **NOT NULL** also optionally can be specified on other attributes.
: For example, employee names, phone may be specified by **non null**; What if new employee's name is unknown?
- Referential Integrity
 - This is specified by **FOREIGN KEY REFERENCES**
 - SQL's default action is to **reject** the operation that violates referential integrity violation
 - SQL also provides **user specified trigger** action

CREATE TABLE : Key/Entity Constraints



ER :


Relation :

EMP (SSN, name, bdate)

Table :

```
CREATE TABLE EMP
( SSN    CHAR(9)      NOT NULL
  name   VARCHAR(10) NOT NULL
  bdate  DATE
  PRIMARY KEY (SSN)
  UNIQUE (name)
)
```

CREATE TABLE : Referential Constraints



EMP (SSN, name, bdate)

DEPT (DNO, name, Mgr_SSN)

```
CREATE TABLE EMP
( SSN CHAR(9) NOT NULL
  name CHAR(10) NOT NULL
  bdate DATE
  PRIMARY KEY (SSN)
  UNIQUE (name) )
```

```
CREATE TABLE DEPT
(DNO INT NOT NULL
 name CHAR(10)
 Mgr_SSN CHAR(9) NOT NULL
 PRIMARY KEY(DNO),
 FOREIGN KEY(Mgr-SSN) REFERENCES EMP(SSN))
```


CREATE TABLE : Default/DOMAIN/CHECK

● DEFAULT

- Default value is included in new tuple if explicit value is not provided
- For example, default manager for new department is;

CREATE TABLE DEPT

Mgr_SSN CHAR(9) **DEFAULT** 999999999

- If no default value is specified, the default DEFAULT value is NULL.

● CHECK

- We can restrict domain to specific values;
- For example, department numbers are between 1 and 20;

DNO **INT** . . **CHECK** (DNO > 0) **AND** (DNO < 21)

- CHECK can be also used with **CREATE DOMAIN**

CREATE DOMAIN DNO **AS INT**

DNO **INT** . . **CHECK** (DNO > 0) **AND** (DNO < 21)

CREATE DOMAIN / DEFAULT

```
CREATE DOMAIN SSN-TYPE AS CHAR(9)
CREATE TABLE EMP
( SSN SSN-TYPE NOT NULL
  name CHAR(10) NOT NULL
  bdate DATE
  PRIMARY KEY (SSN)
  UNIQUE (name) )
```

```
CREATE TABLE DEPT
( DNO INT NOT NULL
  name CHAR(10) NOT NULL
  Mgr-SSN SSN-TYPE
  PRIMARY KEY(DNO),
  UNIQUE(name),
  FOREIGN KEY(Mgr-SSN) REFERENCES EMP(SSN)
```

- What if we change later SSN into another type?

CREATE TABLE : Trigger

- SQL provides the following "Referential Trigger Action";

Action (Referencing)	Event (Referenced)
CASCADE SET NULL SET DEFAULT	ON DELETE ON UPDATE

- Total 6 trigger actions possible; For example:
 - (1) **ON UPDATE CASCADE**
: Change the value of referencing FK to the updated (new) PK value for all the referencing tuples;
 - (2) **ON DELETE CASCADE**
: Delete all the referencing tuples;
 - (3) **ON UPDATE SET DEFAULT**
 - (4) **ON DELETE SET DEFAULT**
: The value of FK in referencing tuples is changed to default value;

CREATE TABLE : Trigger


(5) ON UPDATE SET NULL

(6) ON DELETE SET NULL

: The value of FK in referencing tuples is changed to NULL value;

- Note : This **SET NULL** option is not allowed if the FK in referencing relation is a part of its own primary key.
- It is the responsibility of database designer to specify appropriate trigger actions.
- In general, **CASCADE** is useful for following “relationship” relations.
 - binary relationships (i.e., WORK-ON)
 - weak entity types (i.e., DEPENDENT)
 - multi-valued attributes (i.e., DEPT_LOCATIONS)

CREATE TABLE : Trigger



EMP (SSN, name, DNO)

DEPT (DNO, name, Mgr-SSN)

```
CREATE TABLE DEPT
( DNO          INT          NOT NULL
  name         CHAR(10)
  Mgr-SSN     CHAR(9)      NOT NULL
  PRIMARY KEY(DNO),
  FOREIGN KEY(Mgr-SSN) REFERENCES
    EMP(SSN)
    ON DELETE SET NULL
    ON UPDATE CASCADE )
```

```
CREATE TABLE EMP
( SSN CHAR(9) NOT NULL
  name CHAR(10) NOT NULL
  bdate DATE
  PRIMARY KEY (SSN)
  UNIQUE (name) )
```

CREATE TABLE : Trigger

COURSE

CID	name
CS200	OS
CS250	DB
CS300	PL

Referenced:

PK : CID

ENROLL

CID	SID	credit
CS200	12345	3
CS200	23456	3
CS300	23456	4
CS250	23456	3
CS300	45678	3

Referencing:

FK : CID and SID

PK : (CID, SID)

STUDENT


SID	name	age
12345	Bob	22
23456	An	18
34567	Jim	30
45678	Eve	27

Referenced:

PK : SID


- **CASCADE ON DELETE STUDENT**
- **CASCADE ON UPDATE COURSE**
- **SET NULL ON DELETE COURSE** : This is not allowed; Why??

CREATE TABLE : Exercise



- Convert the above ER schema into relational schema;
- Create tables by using SQL; (Specify all constraints/triggers)

CREATE TABLE



EMP (SSN, name, bdate, DNO, Super-SSN)

DEPT (DNO, name, Mgr-SSN)

CREATE TABLE : Exercise

● **CREATE TABLE** DEPT
(DNO INTEGER NOT NULL
name VARCHAR(10)
Mgr-SSN CHAR(9) NOT NULL
PRIMARY KEY (DNO),
UNIQUE (name),
FOREIGN KEY (MgrSSN) **REFERENCES** EMP(SSN)
ON DELETE SET DEFAULT
ON UPDATE CASCADE)

CREATE TABLE EMP
(SSN CHAR(9) NOT NULL
name CHAR(10) NOT NULL
bdate DATE
DNO INTEGER
Super-SSN CHAR(9)
PRIMARY KEY (SSN)
FOREIGN KEY (DNO) **REFERENCES** DEPT(DNO)
ON DELETE SET DEFAULT
ON UPDATE CASCADE,
FOREIGN KEY (Super-SSN) **REFERENCES** EMP(SSN)
ON DELETE SET NULL
ON UPDATE CASCADE)

DROP TABLE

- Used to remove a relation (base table) *and its definition*
- The dropped table can no longer be used in queries, updates, or any other commands since it does no longer exist.
- Two DROP options: **CASCADE** and **RESTRICT**
- Example:

DROP TABLE DEPENDENT **RESTRICT**

- This table is removed only if it is not referenced in any constraints or view. Otherwise, the DROP command will not be executed.
- For example, by foreign keys in another tables or views

DROP TABLE DEPENDENT **CASCADE**

- All constraints and views that reference this table are also removed automatically

ALTER TABLE

- Used to add or drop columns and change column definition to an existing table.
- Example:

ALTER TABLE EMP ADD COLUMN hobby CHAR(12)

- The new column 'hobby' will have NULLs in all tuples if no default value is specified.
- Users must enter a value for 'hobby' column by using **UPDATE** command.

ALTER TABLE EMP DROP COLUMN bdate CASCADE

- All constraints and views that 'bdate' column are also dropped automatically.
- If **RESTRICT** is used, this command is executed only if no views or constraints reference 'bdate' column.

Retrieval: Queries

- SQL has one basic statement for retrieving information from a database; the **SELECT** command.
- **SELECT** is not the same as the $\text{SELECT}(\sigma)$ operation of the relational algebra
- Important difference between SQL and relational model;
 - SQL allows a table to have two or more duplicate tuples
 - Thus, an SQL relation (table) is a *multi-set* of tuples (= called a **bag**); it is not a set of tuples.
- SQL relations can be constrained to be a **set** by specifying:
 - **PRIMARY KEY** or **UNIQUE** attributes, or
 - **DISTINCT** option in a query

Retrieval Queries

- Basic form of the Retrieval Queries:

```
SELECT  <attribute list>  
FROM    <table list>  
WHERE   <condition>
```

- <attribute list> is a list of attribute names whose values are to be retrieved by the query
- <table list> is a list of the relation names to process query
- <condition> is a conditional (Boolean) expression that identifies the tuples to be retrieved by the query

Figure 3.6

One possible database state for the COMPANY relational database schema.

EMPLOYEE

Fname	Minit	Lname	<u>Ssn</u>	Bdate	Address	Sex	Salary	Super_ssn	Dno
John	B	Smith	123456789	1965-01-09	731 Fondren, Houston, TX	M	30000	333445555	5
Franklin	T	Wong	333445555	1955-12-08	638 Voss, Houston, TX	M	40000	888665555	5
Alicia	J	Zelaya	999887777	1968-01-19	3321 Castle, Spring, TX	F	25000	987654321	4
Jennifer	S	Wallace	987654321	1941-06-20	291 Berry, Bellaire, TX	F	43000	888665555	4
Ramesh	K	Narayan	666884444	1962-09-15	975 Fire Oak, Humble, TX	M	38000	333445555	5
Joyce	A	English	453453453	1972-07-31	5631 Rice, Houston, TX	F	25000	333445555	5
Ahmad	V	Jabbar	987987987	1969-03-29	980 Dallas, Houston, TX	M	25000	987654321	4
James	E	Borg	888665555	1937-11-10	450 Stone, Houston, TX	M	55000	NULL	1

DEPARTMENT

Dname	<u>Dnumber</u>	Mgr_ssn	Mgr_start_date
Research	5	333445555	1988-05-22
Administration	4	987654321	1995-01-01
Headquarters	1	888665555	1981-06-19

DEPT_LOCATIONS

<u>Dnumber</u>	<u>Dlocation</u>
1	Houston
4	Stafford
5	Bellaire
5	Sugarland
5	Houston

WORKS_ON

<u>Essn</u>	<u>Pno</u>	Hours
123456789	1	32.5
123456789	2	7.5
666884444	3	40.0
453453453	1	20.0
453453453	2	20.0
333445555	2	10.0
333445555	3	10.0
333445555	10	10.0
333445555	20	10.0
999887777	30	30.0
999887777	10	10.0
987987987	10	35.0
987987987	30	5.0
987654321	30	20.0
987654321	20	15.0
888665555	20	NULL

PROJECT

<u>Pname</u>	<u>Pnumber</u>	Plocation	Dnum
ProductX	1	Bellaire	5
ProductY	2	Sugarland	5
ProductZ	3	Houston	5
Computerization	10	Stafford	4
Reorganization	20	Houston	1
Newbenefits	30	Stafford	4

DEPENDENT

<u>Essn</u>	<u>Dependent_name</u>	Sex	Bdate	Relationship
333445555	Alice	F	1986-04-05	Daughter
333445555	Theodore	M	1983-10-25	Son
333445555	Joy	F	1958-05-03	Spouse
987654321	Abner	M	1942-02-28	Spouse
123456789	Michael	M	1988-01-04	Son
123456789	Alice	F	1988-12-30	Daughter
123456789	Elizabeth	F	1967-05-05	Spouse

Simple Query : Single Table

- Retrieve the birthdate and address of employees whose sex is 'male' and whose salary is greater than 30,000.

```
SELECT    bdate, address
FROM      EMP
WHERE     (sex = 'male') AND (salary > 30000)
```

- SELECT**-clause specifies **PROJECT**(π);

WHERE-clause specifies **SELECT** (σ)

(참고: SQL의 SELECT는 relational algebra의 σ 와 다름; π 를 의미)

- This SQL expression is equivalent to:

$\pi_{\text{bdate, address}} (\sigma_{(\text{salary} > 30000) \text{ AND } (\text{sex} = \text{'male'})} \text{EMP})$

Optional WHERE / Use of *

- 'WHERE' is optional; That means "no condition" is applied; Thus, all tuples of the table are retrieved.

```
SELECT  SSN  
FROM    EMP
```

This is equivalent to : $\pi_{SSN}(EMP)$

- To retrieve all the attribute values of the selected tuples, the symbol * is used, which means *all the attributes*

```
SELECT  *  
FROM    EMP  
WHERE   salary > 30000
```

This is equivalent to : $\sigma_{salary > 30000}(EMP)$

SQL Allows Duplicated Tuples

- SQL allows duplicate tuples in a relation : **Multi-Set** of Tuples (= bag)
- Selection ($\sigma_{\theta}(R)$) :
If there are **m** copies of tuple **t** in **R** and **t** satisfies σ_{θ} , then there are **m** copies of **t** in $\sigma_{\theta}(R)$.
- Projection ($\pi_A(R)$) :
For each copy of tuple **t** in **R**, there is a copy of tuple $\pi_A(\mathbf{t})$ in $\pi_A(R)$, where $\pi_A(\mathbf{t})$ denotes the projection of the single tuple **t**.
- Cartesian Product ($R_1 \times R_2$) :
If there are **m1** copies of tuple **t1** in **R1** and **m2** copies of tuple **t2** in **R2**, there are **m1** x **m2** copies of the tuple **t1** * **t2** in **R1** x **R2**

DISTINCT

- SQL does not remove duplicate tuples in query result: Why?
 - Duplicates removal is expensive; (Sort and elimination cost!)
 - Users want to see duplicate tuples in the result of query.
 - When aggregate function (ie., sum, avg) is applied, we need to keep duplicated tuples.
- To remove duplicates in a query result, **DISTINCT** is used.
- Q1 may have duplicate values, but Q2 does not have any duplicates.

Q1 : **SELECT** salary
FROM EMP

Q2: **SELECT** **DISTINCT** salary
FROM EMP

- What do think about the following query Q3?

Q3 : **SELECT** **DISTINCT** SSN, salary
FROM EMP

Multiple Table Query

- If more than one tables are specified in the FROM-clause, then Cartesian Product (X) of tuples is selected.
- Retrieve all the combination of tuples from both employees and departments.

```
SELECT      *  
FROM        EMP, DEPT
```

- This is equivalent : EMP X DEPT

```
SELECT      SSN, DNAME  
FROM        EMP, DEPT
```

- This is equivalent : $\pi_{SSN, DNAME} (EMP \times DEPT)$

Multiple Table Query

- Use a join to query data from more than one table. For example, we have a join from two tables

```
SELECT      <attribute list>
FROM        table1, table2
WHERE        table1.Ai op table2.Bj
```

- Write the join condition in the WHERE clause;
A_i and B_j : Join Attributes,
op ∈ {=, ≠, >, ≥, <, ≤})
- Prefix the attribute name with the table name when the same attribute name appears in more than one table.

Multiple Table Query

- Retrieve SSN and age of employees who work for 'research' dept.

```
SELECT      SSN, age
FROM        EMP, DEPT
WHERE       (DNAME = 'research') AND
              (DNO = DNUMBER)
```

- (DNO = DNUMBER) condition specifies the JOIN;
- This SQL expression is equivalent to :

(1) $\pi_{SSN, age} (\sigma_{(DNAME = 'research')} (EMP \bowtie_{DNO = DNUMBER} DEPT))$
or
(2) $\pi_{SSN, age} ((EMP \bowtie_{DNO = DNUMBER} (\sigma_{(DNAME = 'research')} DEPT)))$

Multiple Table Query

- For project name 'product X', list the project number, controlling department number, and the department manager's salary.

```
SELECT      PNO, DNO, salary
FROM        PROJECT, DEPT, EMP
WHERE       (PNAME = 'product X') AND
              (Mgr_SSN = SSN) AND
              (DNUM = DNUMBER)
```

- Notice that there are *two* join conditions
 - (DNUM = DNUMBER) relates a project to its controlling department
 - (Mgr-SSN = SSN) relates the controlling department to the employee who manages that department

Rename Operation

- We can rename for both relation and attributes, taking the form;
old name **AS** new name
- The **AS** clause can appear in both the **SELECT** and **FROM** clause.

SELECT	SSN, address AS addr
FROM	EMPLOYEE
WHERE	(age >= 50)

SELECT	SSN, address
FROM	EMPLOYEE AS EMP
WHERE	(age >= 50)

Attribute Name Conflicting

- Retrieve name and dept. phone of employees who work for each dept.

EMP(SSN, name, phone, DNO)

DEPT(DNO, name, phone)

```
SELECT    name, phone
FROM      EMP, DEPT
WHERE     DNO = DNO
```

Which name? which phone?
: Name conflicting!!

```
SELECT    EMP.name, DEPT.phone
FROM      EMP, DEPT
WHERE     EMP.DNO = DEPT.DNO
```

방식 1 :
Prefix에 Relation name
을 부쳐 구별

```
SELECT    e.name, d.phone
FROM      EMP AS e, DEPT AS d
WHERE     e.DNO = d.DNO
```

방식 2 :
Tuple variable로 구별
(뒷 장 참조!)

Use of Tuple Variables

- We can use **tuple** variables to refer any **tables** in SQL query;
- Retrieve name and dept. phone of employees who work for 'research' dept.

```
SELECT    e.name, d.phone
FROM      EMP AS e, DEPT AS d
WHERE     (d.name = 'research') AND
           (e.DNO = d.DNUMBER)
```



- The alternative table names **e** and **d** are called "**tuple variables**"
- This SQL expression is **equivalent** to :

```
SELECT    EMP.name, DEPT.phone
FROM      EMP, DEPT
WHERE     (DEPT.name = 'research') AND
           (EMP.DNO = DEPT.DNUMBER)
```

Use of Tuple Variables

- Tuple variables are very useful for comparing two tuples in the same relation.
- In this case, two tuple variables are given to the same relation.
- For each employee, retrieve the employee's SSN, and the name of his (or her) direct supervisor.

```
SELECT      e.SSN, s.name
FROM        EMP AS e, EMP AS s
WHERE       e.Super_SSN = s.SSN
```

- The alternative table names **e** and **s** are called "tuple variables"
- We can think of **e** and **s** as *two different copies* of EMPLOYEE;
 - **e** represents employees in role of 'workers'
 - **s** represents employees in role of 'bosses'

Use of Tuple Variables

- Retrieve SSN and names of all employees whose salary greater than Smith's salary.

```
SELECT      e.SSN, e.name
FROM        EMP AS e, EMP AS s
WHERE       (e.salary > s.salary) AND
            (s.name = 'smith')
```

- We can also think of **e** and **s** as *two different copies* of EMPLOYEE;
 - **e** represents in the role of employees
 - **s** represents in the role of employee(s) with name = 'smith'

Use of Tuple Variables

- Retrieve the names of two employees (i.e., married) who share the same address.

```
SELECT      e1.name, e2.name  
FROM        EMP AS e1, EMP AS e2  
WHERE       (e1.address = e2.address) AND  
            (e1.name < e2.name)
```

- The second condition ($e1.name < e2.name$) means that the name of first employee precedes the name of second employee alphabetically.
- If this condition is omitted, tuple variables $e1$ and $e2$ both could refer to the same tuple. Thus, the query result produces each employee name paired with itself.
- Question: If use " $<>$ " instead of " $<$ ", what will happen?

Substring Pattern Matching

- SQL provides **LIKE** comparison operator for pattern matching. We describe patterns by using two special characters.
 - **Percent(%)**: The % matches any "substring".
 - **Underscore(_)**: The _ matches any "character".
- Patterns are case sensitive; that is, upper case characters do not match lower case characters, or vice versa; Some examples are;
 - 'Blue%' matches any string beginning with 'Blue'.
 - '%idge%' matches any string containing "idge" as substring.
Examples: Blueridge, Ridgeriver, 'Rockbridge, Ridgeway, . .
 - ' _ _ ' matches any string of exactly 3 characters.
 - ' _ _ _%' matches any string of at least 3 characters.

Substring Pattern Matching

- Retrieve all employees whose address is in Houston, Texas.

```
SELECT      *  
FROM        EMP  
WHERE       Address LIKE '%Houston, Texas%'
```

- Retrieve all employees who were born during the 1970s.

```
SELECT      *  
FROM        EMP  
WHERE       Bdate LIKE '__ 7 _ _ _ _ _ _ _'
```

- Note that **NOT LIKE** also can be used in a similar way.

Arithmetic Operations

- We can use **arithmetic operations** in query; The standard operators **+, -, *, and /** can be applied to attributes with numeric data types;
- Show resulting salaries for every employee working for the research department is given 10% raise.

```
SELECT    e.SSN, 1.1 * e.salary AS increased_sal  
FROM      EMP AS e  DEPT d  
WHERE      (d.name = 'research') AND (e.DNO = d.DNUMBER)
```

- SQL provides **BETWEEN** operator to simplify comparison operators. For example, retrieve all employees whose salary is between 30,000 and 40,000. (that is, salary **>= AND** salary **<=** 40000)

```
SELECT      *  
FROM        EMP  
WHERE        salary BETWEEN 30000 AND 40000
```


Set Operators

- **UNION, INTERSECT, EXCEPT**
- Each of the above operations automatically removes duplicated tuples;
- To retain all duplicated tuples, use the corresponding operations **UNION ALL, INTERSECT ALL, EXCEPT ALL**.
- Suppose a tuple **t** occurs **m** times in R and **n** times in S;
 - R **UNION ALL** S : (**m** + **n**) times occurs.
 - R **INTERSECT ALL** S : **Min(m, n)** times occurs.
 - R **EXCEPT ALL** S : **Max(0, m – n)** times occurs.

Set Operators

- Retrieve SSN and name of employees with age > 40 or sex = 'male'.

```
SELECT      SSN, name
FROM        EMP
WHERE       age > 40
UNION
SELECT      SSN, name
FROM        EMP
WHERE       sex = 'male'
```

- This is equivalent to :

```
SELECT      SSN, name
FROM        EMP
WHERE       (age > 40) OR (sex = 'male')
```

Set Operators

- Retrieve SSN of employees with any no dependents.

```
SELECT      SSN
FROM        EMP
EXCEPT
SELECT      ESSN
FROM        DEPENDENT
```

Set Operators

- Retrieve all project numbers for projects that involve an employee whose name is 'Bob' as a worker, and as a manager of the department that controls the project.

```
SELECT    PNO
FROM      PROJECT, WORK-ON, EMP
WHERE      (PNUMBER = PNO) AND (ESSN = SSN)
             AND (name = 'Bob')
```

INTERSECT

```
SELECT    PNO
FROM      PROJECT, DEPT, EMP
WHERE      (DNUM = DNUMBER) AND (Mgr_SSN = SSN)
             AND (name = 'Bob')
```

NULL Values

- NULL represents missing values; It has 3 different meanings :
 - (1) **Unknown** (exists, but unknown):
 - Someone has 'blood-type', but currently is not known.
 - (2) **Unavailable** (exists, but is withheld):
 - Someone has 'phone-number' but does not want it to be listed.
 - (3) **No applicable** (undefined or non-existent):
 - 'spouse' may be NULL for unmarried someone.
- It is not possible to distinguish these 3 meanings: For example, Suppose many persons have "NULL value for phone-number"; How can we interpret this?

NULL Values

- NULL values occur inevitably due to the following cases;
 - Outer Union
 - Outer Join
 - Insertion of tuples with unknown attributes
- Let X have the value NULL;
 - (1) Any arithmetic operation with NULL returns "**NULL**".
 - $X + 5$, $X - 5$, $X * 5$, . . .
 - (2) Any comparison with NULL returns "**UNKOWN**" :
 - $X < 5$, $X = 5$, $X = \text{NULL}$, $X <> 5$, . . .
- Note that "**UNKOWN**" is another truth value like TRUE and FALSE.
- SQL provides **IS NULL** (or **IS NOT NULL**) for dealing NULL values; "X IS NULL" returns "TRUE" if X has NULL value; It returns "FALSE" otherwise.

NULL Values

- SQL uses a **three**-valued logic: TRUE, FALSE, UNKNOWN

AND	T	F	U
T	T	F	U
F	F	F	F
U	U	F	U

OR	T	F	U
T	T	T	T
F	F	F	U
U	U	U	U

	NOT
T	F
F	T
U	U

- To understand how 3-valued logic works:
Let $T = 1$, $F = 0$, $U = \frac{1}{2}$,
 $AND = MIN$, $OR = MAX$, $NOT(X) = 1 - X$.

- Example:

$$\begin{aligned} T \text{ AND } (F \text{ OR } NOT(U)) &= MIN(1, MAX(0, (1 - \frac{1}{2}))) \\ &= MIN(1, MAX(0, \frac{1}{2})) = MIN(1, \frac{1}{2}) = \frac{1}{2} (= U). \end{aligned}$$

NULL Values

- Result of **WHERE** condition is treated as **FALSE** if it evaluates to **UNKNOWN**.
- Tuples that are evaluated to only **TRUE** are selected.
 - Thus, tuples with either **FALSE** or **UNKNOWN** values are not selected from the result.

- Example:

```
SELECT SSN
FROM EMP
WHERE (age < 50) AND
      ((sex = 'male') OR (salary > 50000))
```

SSN	age	sex	salary
22222	null	null	60000

- Is this tuple selected? Answer: NO!

NULL Values

- (Funny) Example:

```
SELECT SSN
FROM EMP
WHERE (salary < 50000) OR (salary >= 50000)
```

SSN	age	sex	salary
33333	40	male	null

- Is this tuple selected? Answer: NO!

- In this case, we need to use **IS NULL** to check whether the salary attribute has NULL or not.

```
SELECT SSN
FROM EMP
WHERE (salary < 50000) OR (salary >= 50000)
OR (salary IS NULL)
```

- Now, all tuples are selected!

Use of JOIN Operators

- SQL permits users to specify join operators in **FROM** clause; This is easier to understand than mixing together select and join conditions in **WHERE**;

```
SELECT  SSN, age  
FROM    EMP JOIN DEPT ON DNO = DNUMBER  
WHERE   DNAME = 'research'
```

This is equivalent to:

```
SELECT  SSN, age  
FROM    EMP, DEPT  
WHERE   (DNAME = 'research') AND (DNO = DNUMBER)
```

- SQL also provides various types of joins;
 - **JOIN, NATURAL JOIN, LEFT OUTER JOIN, RIGHT OUTER JOIN, CROSS JOIN, . . .**

Use of JOIN Operators

- In the case of **Natural Join**, if the name of join attributes are not the same, it is possible to rename the attributes so that they match.
- Consider the following query;

```
SELECT      Name, age
FROM        EMP NATURAL JOIN DEPARTMENT
              AS DEPT(Dname, DNO, MgrSSN, Mdate)
WHERE       Dname = 'research'
```

- Not that DEPARTMENT is renamed as DEPT and Dnumber is also renamed as DNO to match the desired join attribute DNO in EMP relation.

Use of JOIN Operators

- In this join, only employees who have a supervisor are included in query result;

```
SELECT      e.SSN, s.name
FROM        EMP AS e, EMP AS s
WHERE       e.Super-SSN = s.SSN
```

- To include all employees in query result, the following outer join can be used;

```
SELECT      e.SSN, s.name
FROM        EMP AS e LEFT OUTER JOIN EMP AS s
ON e.Super-SSN = s.SSN
```

- It is also possible to nest many join operators;

```
SELECT      SSN, age
FROM        (PROJECT JOIN WORK-ON ON PNO = PNUMBER)
              (EMP JOIN WORK-ON ON SSN = ESSN) )
WHERE       PNAME = 'laptop'
```

Nested Queries (= Subqueries)

- A query that is part of another query is called a **subquery** (or **nested query**);
 - Subquery can have many nested subqueries as we want;
 - There are many other ways that subquery can be used;
- (1) Subquery can return **constant value**(s), and this value can be compared with another value in WHERE clause;
 - (2) Subquery can return **relation** (a set of tuples) that can be used in various ways in WHERE clause;
 - (3) Subquery can have their relations appear in **FROM** clause;

Nested Queries

- A **nested query** can be specified within the **WHERE**-clause of another query, called the **outer query**

- Single query :

```
SELECT    SSN
FROM      EMP, DEPT
WHERE     (DNAME = 'research') AND DNO = DNUM
```

- Nested Query : Having executed nested query returns 'research' number; Then, outer query returns SSN(s) only if its DNO is equal to result of nested query.

```
SELECT    SSN
FROM      EMP
WHERE     DNO =
```

Outer query : Select a EMP tuple if its DNO is equal to result of nested query

```
(SELECT    DNUM
FROM      DEPT
WHERE     DNAME = 'research')
```

Nested query :
research dept number

- Question: What if there are more than one research dept numbers?

Nested Queries : IN

- A **nested query** can be compared with **outer query** by using **IN** operator;
- (**v IN V**) compares a value **v** (in outer query) with the value(s) **V** (in nested query), and returns "true" if **v** is **included in V**. Otherwise, return "false";

```
SELECT SSN  
FROM EMP  
WHERE DNO
```

Outer query : Select an EMP tuple if its DNO is included in result of nested query)

IN

```
SELECT DNUM  
FROM DEPT  
WHERE DNAME = 'research'
```

Nested query :
research dept number(s)

Nested Queries : IN

- SQL allows use of (sub)**tuples** of values in comparisons by placing them within parenthesis;
- Get SSNs of employees who work for the same (project, hours) on some project that employee with SSN = '123456789' works on;

```
SELECT  DISTINCT ESSN  
FROM    WORK-ON  
WHERE   (PNO, hours)
```

IN

```
SELECT  PNO, hours  
FROM    WORK-ON  
WHERE   ESSN = '123456789'
```

outer query : Select an EMP tuple if its (pno, hours) is included in result of nested query)

nested query :
a set of (pno, hours)
worked by '123456789'

Nested Queries : IN

- Get names of employees who work for the same (project, hours) on some project that employee with SSN = '123456789' works on;

```
SELECT  ename  
FROM    EMP  
WHERE   SSN
```

nested query : Select an EMP tuple name if its SSN is included in result of nested query

IN

```
SELECT  ESSN  
FROM    WORK-ON  
WHERE   (PNO, hours)
```

nested query : Select an EMP ESSN if its (pno, hours) is included in result of nested query)

IN

```
SELECT  PNO, hours  
FROM    WORK-ON  
WHERE   ESSN = '123456789'
```

nested query :
a set of (pno, hours)
worked by '1234567'

Nested Queries : NOT IN

- (**v NOT IN V**) compares a value **v** (in outer query) with the value(s) **V** (in nested query), and returns "true" if **v** is not included in **V**. Otherwise, return "false";
- Retrieve names of employees with no dependents.

```
SELECT  ename
FROM    EMP
WHERE   SSN
```

Outer query : Select an EMP tuple if its SSN is not included in result of nested query)

NOT IN

```
SELECT  ESSN
FROM    DEPENDENT
```

Nested query :
SSN(s) with dependents

- Retrieve names of employees whose names are neither Smith nor Jones.

```
SELECT  ename
FROM    EMP
WHERE   ename
        NOT IN ('Smith', 'Jone')
```

Nested Queries : Use of Tuple Variables

- Nested Queries also can use tuple variables;
- Retrieve names of employees who have a dependent with the same sex as the employees.

```
SELECT    name
FROM      EMP AS e
WHERE     e.SSN
```

outer query : Select a EMP tuple if its SSN is included in result of nested query.

IN

```
SELECT    d.ESSN
FROM      DEPENDENT AS d
WHERE     e.sex = d.sex
```

nested query : Set of employee SSNs with the same sex as their dependents;

Nested Queries : SOME (= ANY)

- **SOME** (= **ANY**) can be used with {<, >, =, <=, >=, <>} operators:
 - For example, (**v** > **SOME** **V**) returns "true" if the value **v** is greater than some of the values in a set **V**.
- Get SSNs of employees whose salary is greater than salary of some employees in department 5:

```
SELECT  SSN  
FROM    EMP  
WHERE   salary
```

> **SOME**

```
SELECT  salary  
FROM    EMP  
WHERE   DNO = 5
```

outer query : Select an EMP tuple if its salary is greater than some values in result of nested query

nested query : a set of all salaries of employees who work for DNO = 5

Example : Meaning of SOME

- Example :

- ✓ $(5 < \text{SOME } \{0, 5, 7\}) = \text{true}$

- ✓ $(5 < \text{SOME } \{1, 3, 5\}) = \text{false}$

- ✓ $(5 = \text{SOME } \{0, 5, 7\}) = \text{true}$

- ✓ $(5 <> \text{SOME } \{0, 5, 7\}) = \text{true}$

- 참조: $(= \text{SOME})$ 은 **IN** 과 동일한 표현임. 반면에, $(<> \text{SOME})$ 은 **NOT IN** 과 동일하지 않음.

Nested Queries : ALL

- **ALL** can be used with {<, >, =, <=, >=, <>} operators:
 - For example, ($v > \text{ALL } V$) returns "true" if the value v is greater than all the values in a set V .
- Get SSNs of employees whose salary is greater than salary of all the employees in department 5:

```
SELECT    SSN  
FROM      EMP  
WHERE     salary
```

> ALL

```
SELECT    salary  
FROM      EMP  
WHERE     DNO = 5
```

outer query : Select an EMP tuple if its salary is greater than all values in result of nested query

nested query : a set of all salaries of employees who work for DNO = 5

Example : Meaning of ALL

- Example :

- ✓ $(5 < \mathbf{ALL} \{0, 5, 6\}) = \text{false}$

- ✓ $(5 > \mathbf{ALL} \{0, 3, 4\}) = \text{true}$

- ✓ $(5 = \mathbf{ALL} \{4, 5, 7\}) = \text{false}$

- ✓ $(5 <> \mathbf{ALL} \{4, 6, 8\}) = \text{true}$

- 참조: ($<> \mathbf{ALL}$)은 **NOT IN** 과 동일한 표현임. 반면에, ($= \mathbf{ALL}$)은 **IN** 과 동일하지 않음.

Nested Queries : EXISTS

- **EXISTS**(Q) returns "true" if there is at least one tuple (= **not empty**) in the result of a nested query Q. Otherwise, it returns "false"
- Retrieve SSNs of employees who have dependent(s).

```
SELECT SSN  
FROM EMP  
WHERE
```

outer query : For each employee, if result of its related nested query is not empty, select that employee SSN.

EXISTS

```
SELECT *  
FROM DEPENDENT  
WHERE SSN = ESSN
```

nested query : For each EMP tuple, select all dependent tuples whose ESSN matches SSN.

Nested Queries : EXISTS

- Retrieve names of employees who have a dependent with the same sex as the employees.

```
SELECT    name
FROM      EMP AS e
WHERE
```

EXISTS

```
( SELECT    *
  FROM      DEPENDENT AS d
 WHERE      (e.SSN = d.ESSN)
            AND (e.sex = d.sex) )
```

← **outer query** : For each employee tuple, if result of its related nested query is not empty, select that employee name.

← **nested query** : Set of dependent tuples with the same sex as their employees;

Nested Queries : NOT EXISTS

- **NOT EXISTS(Q)** returns "true" if there **no tuple** (= **empty**) in the result of a nested query Q. Otherwise, it returns "false".
- Retrieve SSNs of employees who have no dependent.

```
SELECT  SSN  
FROM    EMP  
WHERE
```

outer query : For each employee tuple, if result of its related nested query is empty, select that employee SSN.

NOT EXISTS

```
SELECT      *  
FROM        DEPENDENT  
WHERE       SSN = ESSN
```

nested query : For each EMP tuple, select all dependent tuples whose ESSN matches SSN.

Nested Queries : EXISTS

- Retrieve names of employees whose sex are not the same with any dependent or with no dependent.

```
SELECT    name
FROM      EMP AS e
WHERE
```

NOT EXISTS

```
( SELECT    *
  FROM      DEPENDENT AS d
 WHERE      (e.SSN = d.ESSN)
            AND (e.sex = d.sex) )
```

← **outer query** : For each employee tuple, if result of its related nested query is not empty, select that employee name.

← **nested query** : Set of dependent tuples with the same sex as their employees;

Nested Queries : NOT EXISTS

- Retrieve names of employees whose salary are greater than salary of all the employees with age < 25.

```
SELECT    e1.ename  
FROM      EMP AS e1  
WHERE
```

NOT EXISTS

```
SELECT      *  
FROM        EMP AS e2  
WHERE      e2.age < 25 AND  
            e1.salary <= e2.salary
```

- Question: Suppose that age of all employees are ≥ 25 ; Then, what is the query result?

Division Query : NOT EXISTS and EXCEPT

- Retrieve name of each employee who work on all the projects controlled by department 5.

```
SELECT    name
FROM      EMP
WHERE
```

NOT EXISTS

```
SELECT    Pnumber
FROM      PROJECT
WHERE     DUM = 5
```

EXCEPT

```
SELECT    PNO
FROM      WORK-ON
WHERE     SSN = ESSN
```

outer query : For each employee, if query 1 – query 2 is empty (that means, that employee works all projects by dept 5) then the answer is selected

nested query 1 : Select all project numbers controlled by dept 5

nested query 2 : Select all project numbers performed currently by each employee

EXISTS and NOT EXISTS

- **INTERSECT vs EXISTS**

```
( SELECT  R.A, R.B  
  FROM    R )  
INTERSECT  
( SELECT  S.A, S.B  
  FROM    S )
```

=

```
SELECT  R.A, R.B  
FROM    R  
WHERE  
  EXISTS  
    ( SELECT  *  
      FROM    S  
      WHERE   R.A=S.A  
              AND R.B=S.B )
```

- **EXCEPT vs NOT EXISTS**

```
( SELECT  R.A, R.B  
  FROM    R )  
EXCEPT  
( SELECT  S.A, S.B  
  FROM    S )
```

=

```
SELECT  R.A, R.B  
FROM    R  
WHERE  
  NOT EXISTS  
    ( SELECT  *  
      FROM    S  
      WHERE   R.A=S.A  
              AND R.B=S.B )
```