

Search II


3. Informed Search Algorithms

Outline


- Recap – uninformed search
- Best-first search
 - ▶ Greedy search
 - ▶ A* search
- Admissible heuristics

Recap - Uninformed Search


BFS – expand **shallowest** node in frontier



UCS – expand node with **lowest cost** so far



DFS – expand **deepest** node in frontier



IDS – DFS in BFS manner


Limit = 0




Limit = 1



Limit = 2



Limit = 3



A strategy is defined by picking the **order of node expansion** (based on **depth level** of node, or **cost** of getting to node (UCS))

Informed Search Strategies


- Uninformed searches keep track of where they are now relative to where they have started (root node) in terms of step-costs
- UCS is good because it is complete and optimal but explores options in “every” direction, no information about goal location
- Need an indication if where we are right now is a good place to be for expanding nodes or a bad place to be for expanding nodes
- Use **problem-specific/domain knowledge** beyond the definition of the problem itself – can find solutions more efficiently than uninformed search
- Normally done using **heuristics**

Heuristics

- A heuristic is a function that *estimates* how close a state is to a goal. Usually designed for a particular search problem (problem-specific)
- Most informed search algorithms include a **heuristic function, $h(n)$** as a component of its evaluation function:

$h(n)$ = estimated cost of the cheapest path from the state at node n to a goal state

- Examples are **Euclidean** distance (straight line distance) and **Manhattan** distance (good for horizontal and vertical movements) for pathing problem



Best-first Search

- General search strategy approach for informed search is called best-first search – idea is to use an **evaluation function, $f(n)$** that estimates the total cost. Node with the *lowest* $f(n)$ is expanded first (most **desirable** node)
- Implementation is similar to uniform-cost search, except for the use of f (total cost) instead of g (**cost so far**) to order the priority queue
- Special cases:
 - Greedy best-first search
 - A* search

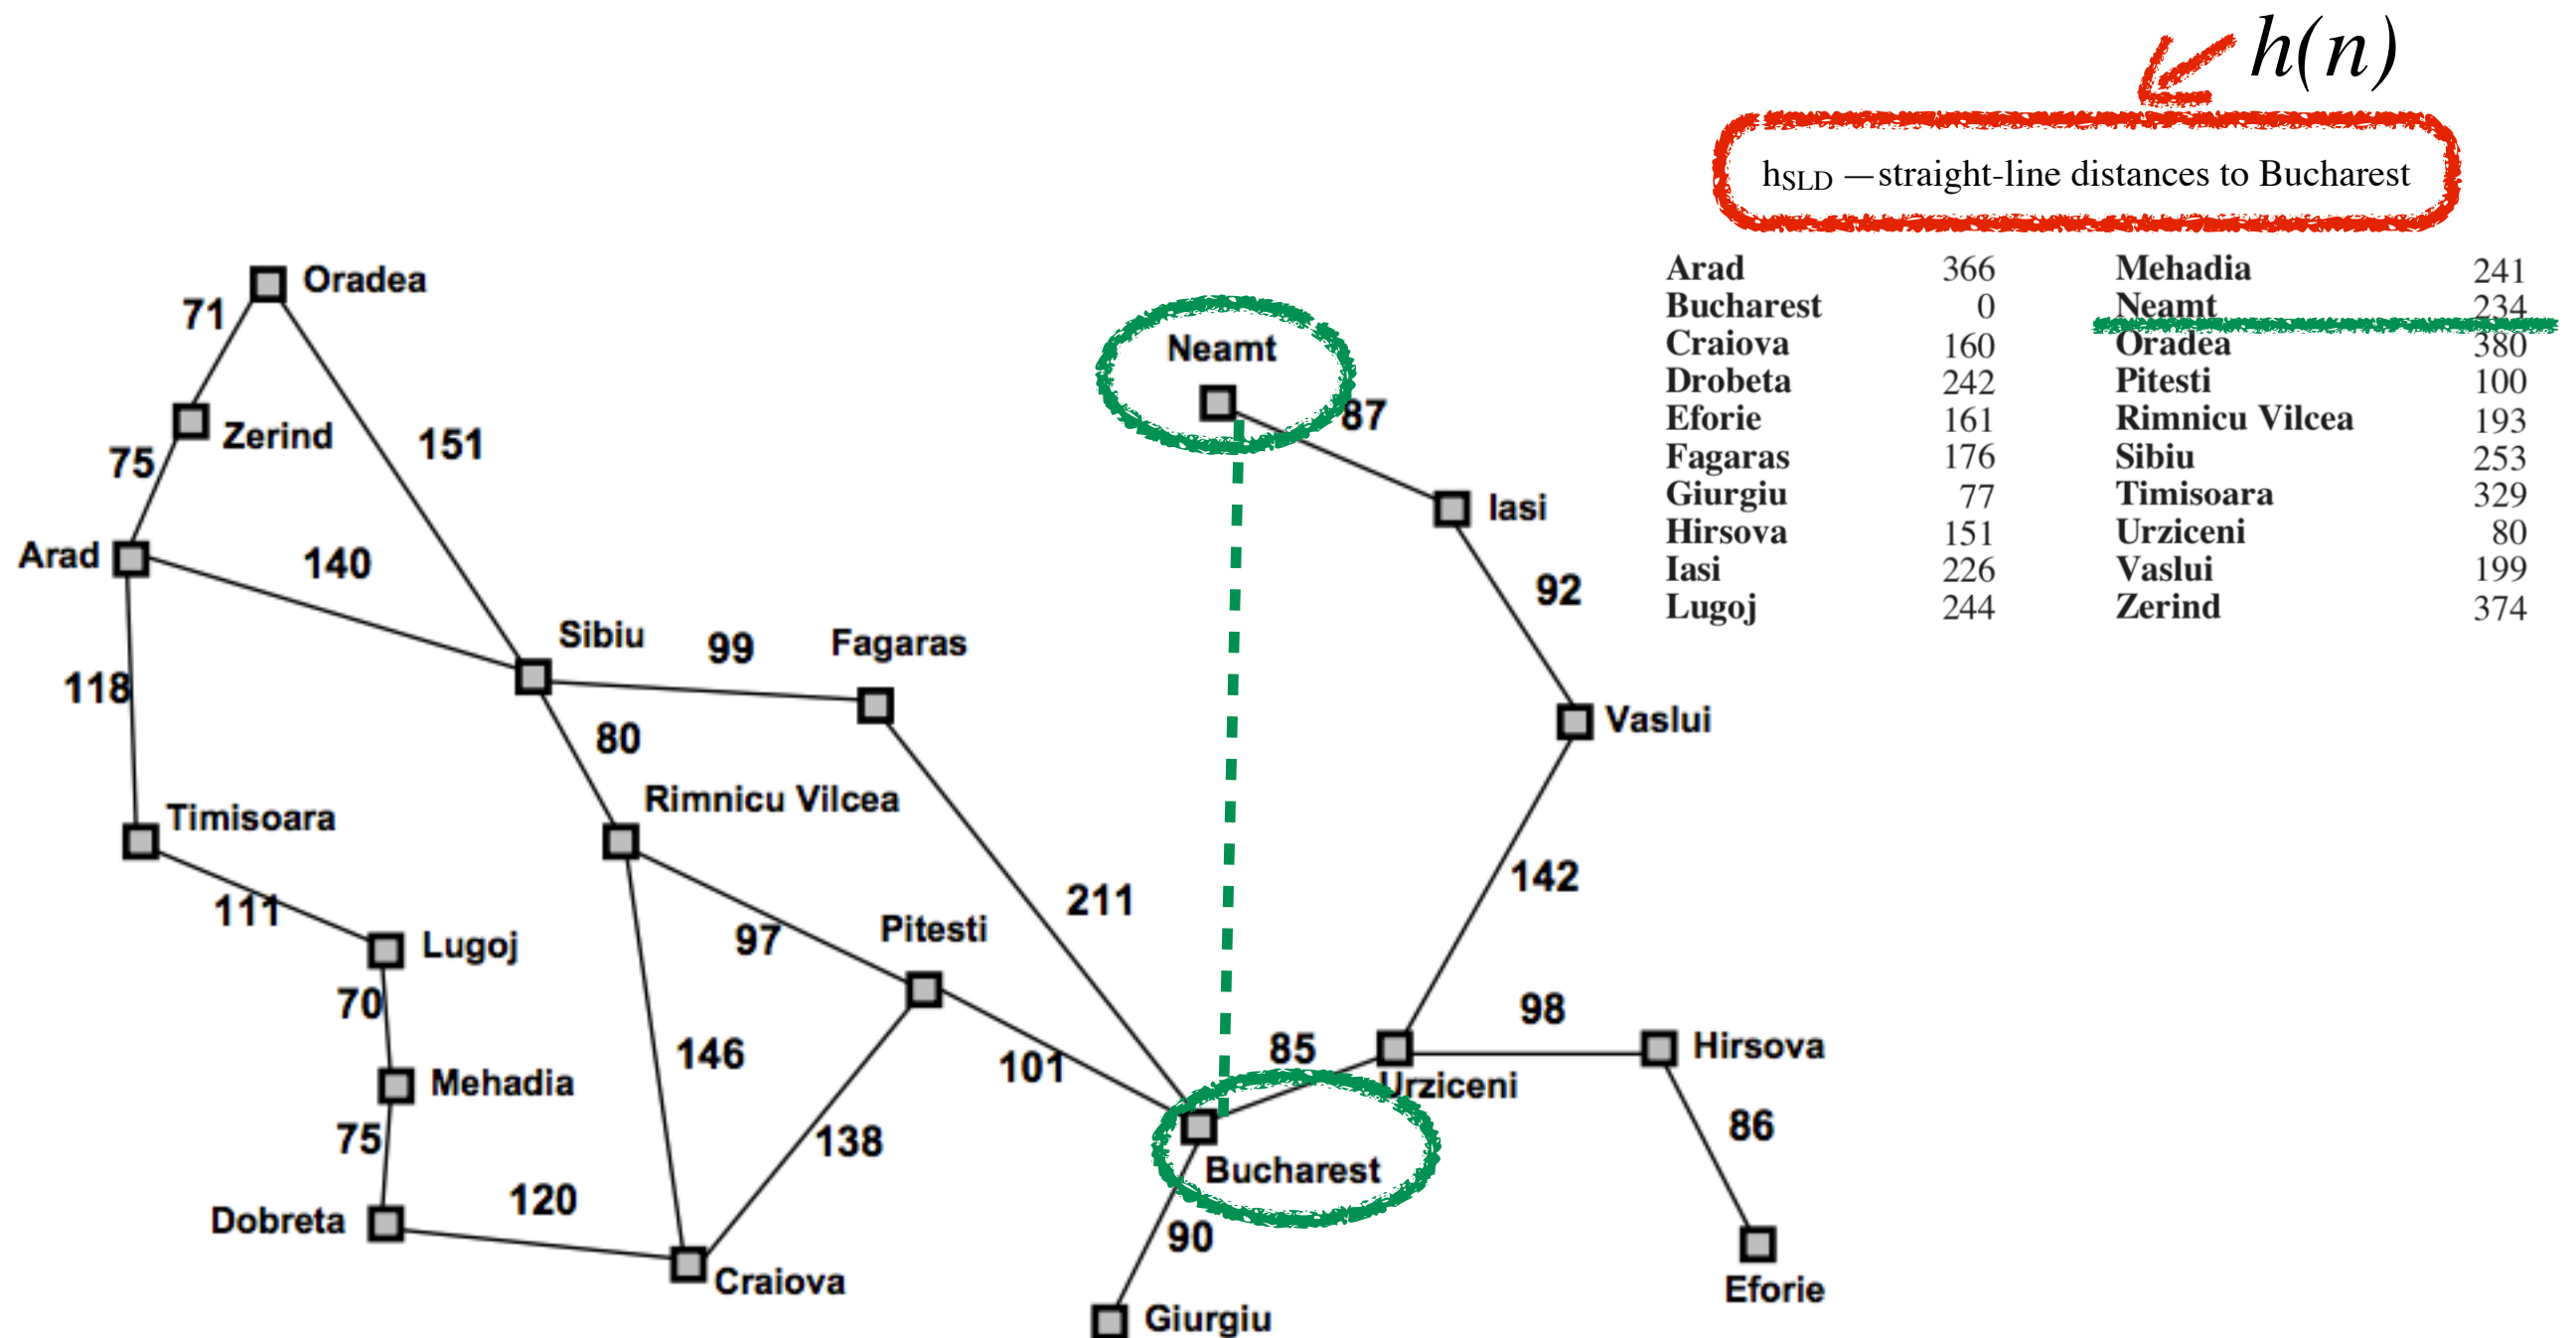
Greedy Best-first Search

- As its name suggests, it tries to expand the node that **appears to be closest** to the goal in order to find the solution quickly
- Evaluates nodes using just the heuristic function:

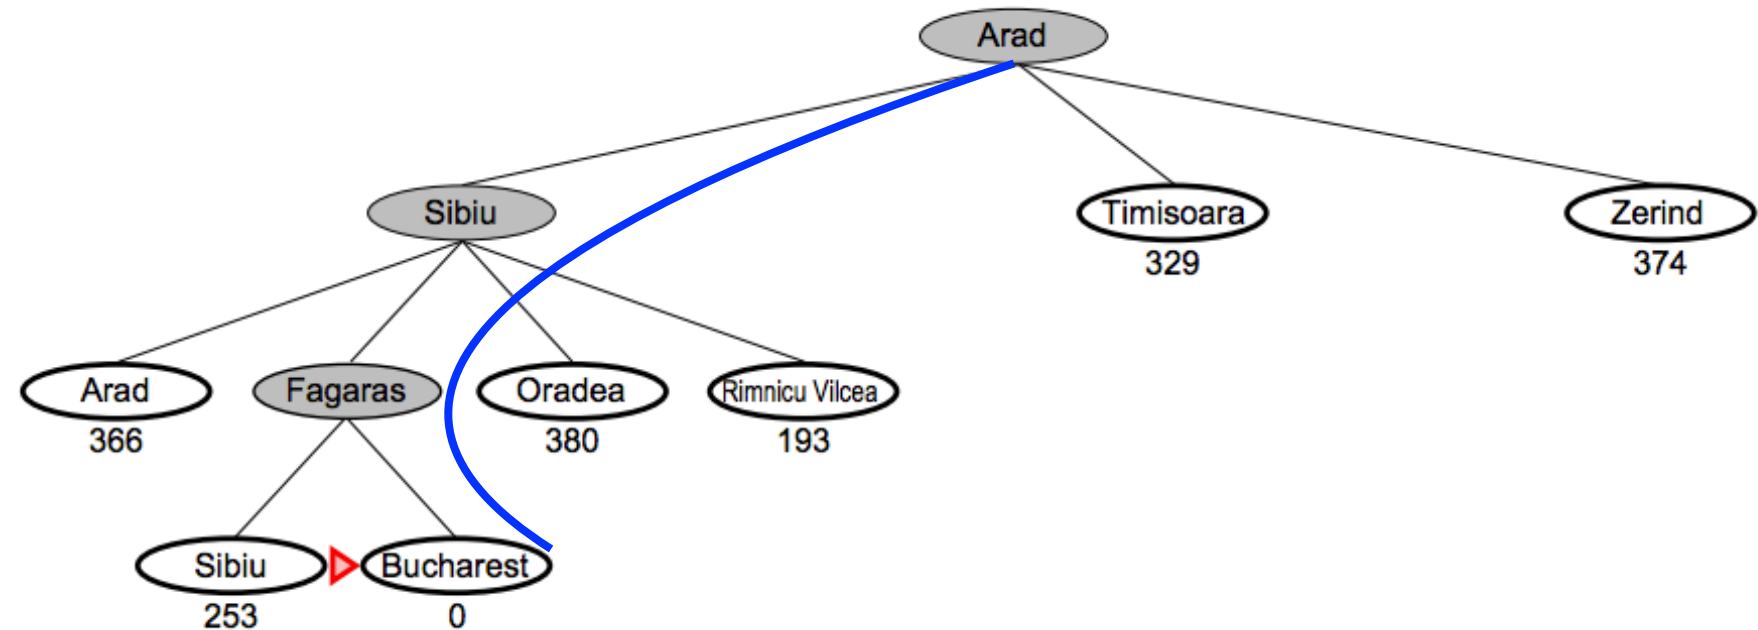



$$f(n) = h(n)$$

- Constraint: goal node will have $h(n) = 0$

Romania with Step Costs (km)

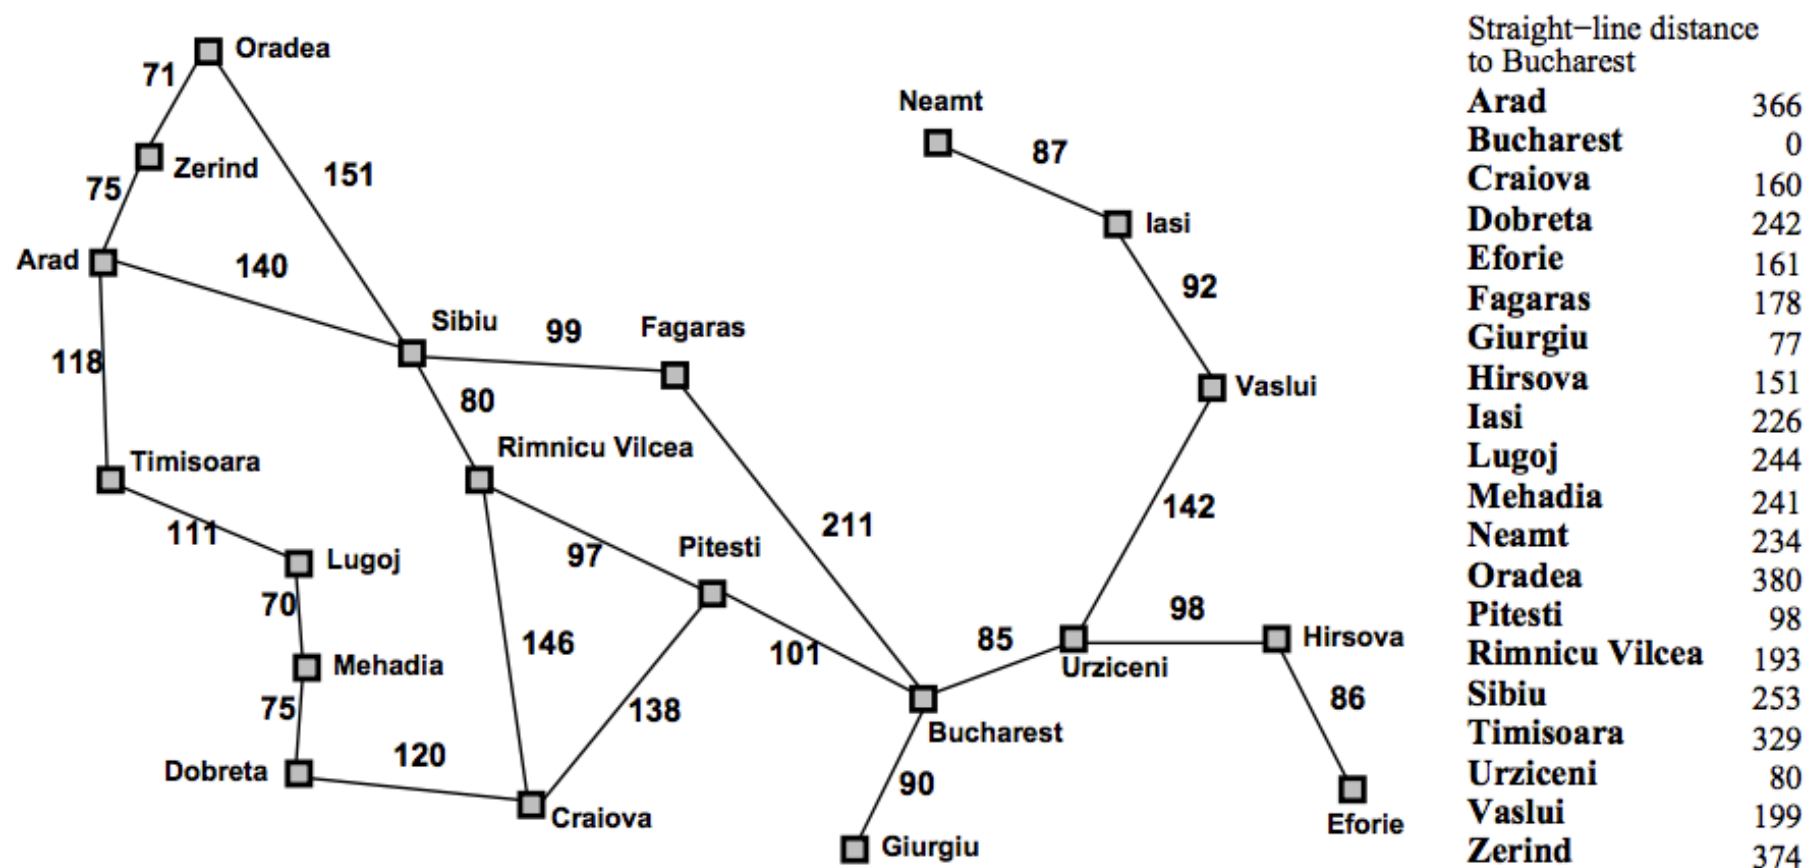


Greedy Search – Arad to Bucharest



Arad	366	Mehadia	241
Bucharest	0	Neamt	234
Craiova	160	Oradea	380
Drobeta	242	Pitesti	100
Eforie	161	Rimnicu Vilcea	193
Fagaras	176	Sibiu	253
Giurgiu	77	Timisoara	329
Hirsova	151	Urziceni	80
Iasi	226	Vaslui	199
Lugoj	244	Zerind	374

Greedy Search – Properties



- **Complete?** No, can get stuck in loops, e.g., *Iasi-Neamt-Iasi-Neamt- ...* (tree-search version)
- **Optimal?** No. In this example, *Arad-Sibiu-Fagaras-Bucharest* is 32kms longer than *Arad-Sibiu-Rimnicu Vilcea-Pitesti-Bucharest*
- **Time?** $O(b^m)$ but a good heuristic can give dramatic improvement
- **Space?** $O(b^m)$ – keeps all nodes in memory

A* Search

- Idea: avoid expanding paths that are already expensive
- Combines UCS with Greedy
- Evaluates nodes by combining $g(n)$, the **cost so far to reach node n** , and $h(n)$, which is the **estimated cost** to goal from n

$$f(n) = g(n) + h(n)$$

$f(n)$ = estimated cost of the cheapest solution through n

- Want $f(n)$ to be minimal
- If $h(n)$ satisfies some conditions, A* search is both complete and optimal
- Identical to UCS except that A* uses $g+h$ instead of g

A* Tree Search Pseudocode

A* tree search: $f(n) = g(n) + h(n)$

```
function A-STAR-SEARCH returns a solution or failure
  node  $\leftarrow$  INITIAL-STATE
  frontier  $\leftarrow$  a priority queue ordered by  $f(n) = g(n) + h(n)$ , with node as the only element
  loop do
    if EMPTY(frontier) then return failure
    node  $\leftarrow$  POP(frontier) /* chooses the node with lowest  $f(n)$  */
    if node == Goal node then return SOLUTION(node)
    for each action in ACTIONS(node.STATE) do
      child  $\leftarrow$  CHILD-NODE(node, action)
      frontier  $\leftarrow$  INSERT(child, frontier)
```

A* Search – Arad to Bucharest


(a) The initial state



frontier	g	h	f
[A]	0	366	366

1


(b) After expanding Arad



[A-S]	140	253	393
[A-T]	118	329	447
[A-Z]	75	374	449

2

(c) After expanding Sibiu

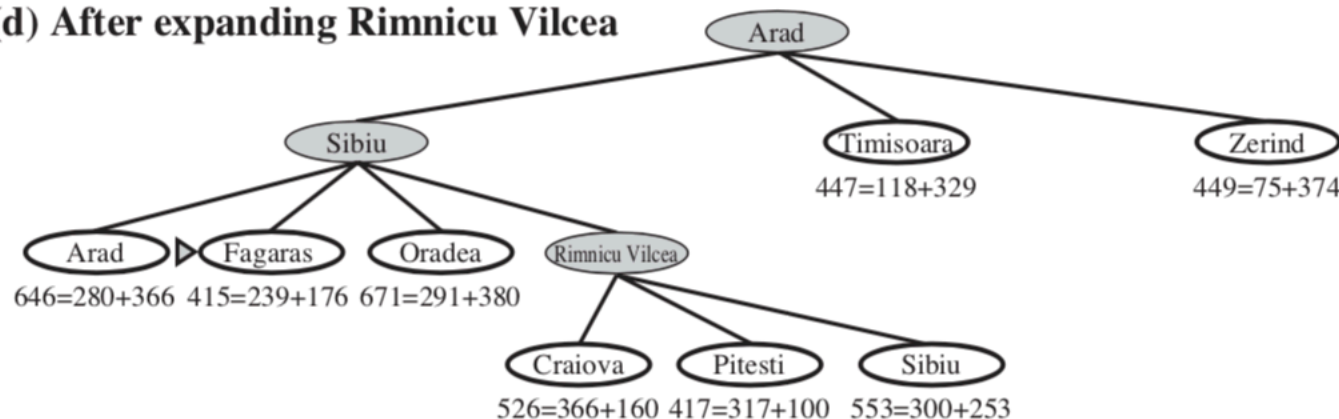


[A-S-A]	280	366	646
[A-S-F]	239	176	415
[A-S-O]	291	380	671
[A-S-RV]	220	193	413

4

3

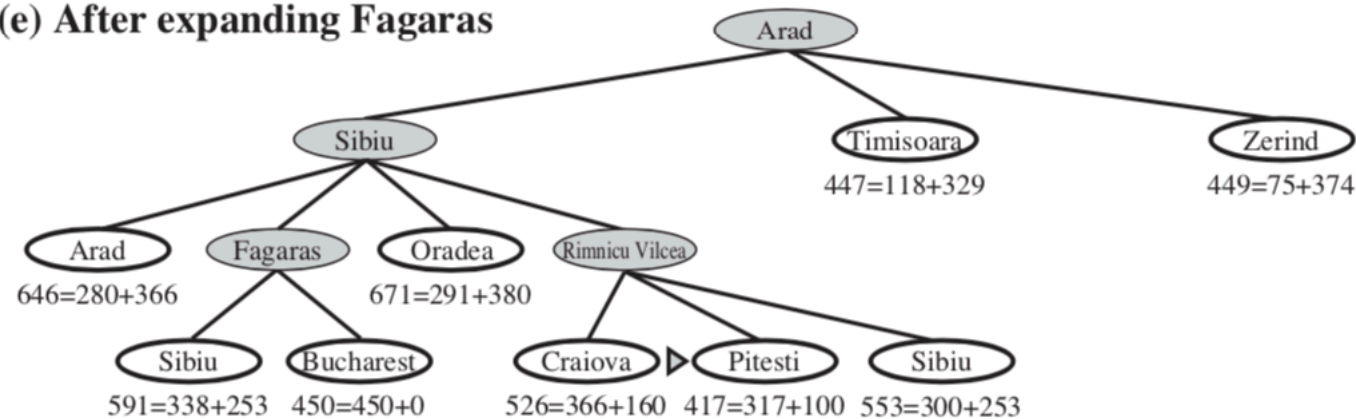
(d) After expanding Rimnicu Vilcea



[A-S-RV-C]	366	160	526
[A-S-RV-P]	317	100	417
[A-S-RV-S]	300	253	553

A* Search – Arad to Bucharest

(e) After expanding Fagaras



frontier	g	h	f
[A-T]	118	329	447
[A-Z]	75	374	449

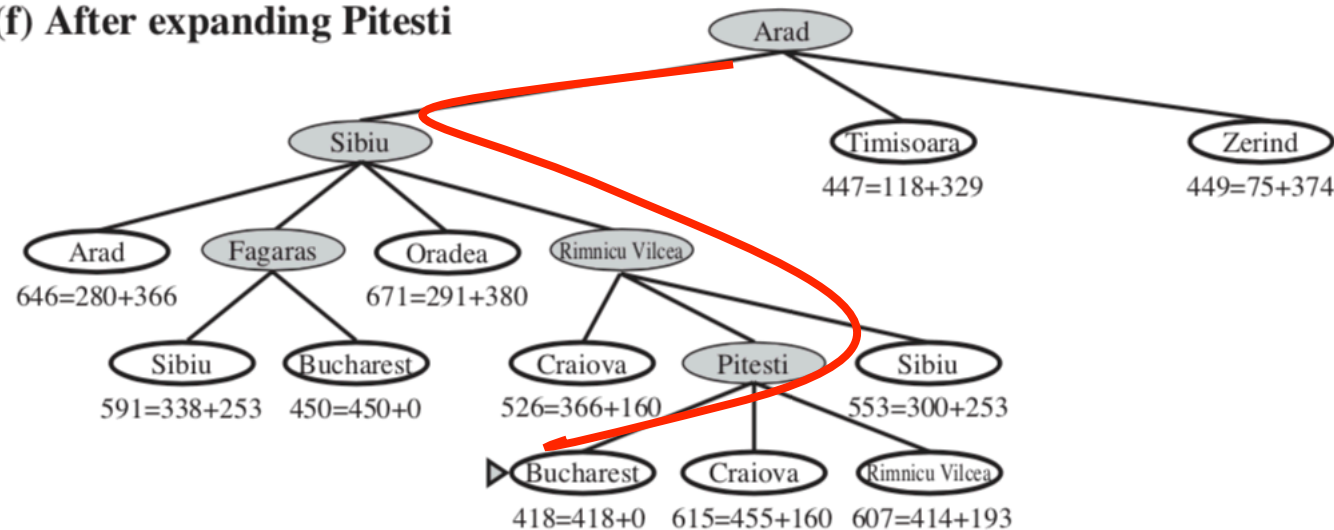
[A-S-A]	280	366	646
[A-S-F]	239	176	415
[A-S-O]	291	380	671

4

[A-S-RV-C]	366	160	526
[A-S-RV-P]	317	100	417
[A-S-RV-S]	300	253	553

5

(f) After expanding Pitesti





[A-S-F-S]	338	253	591
[A-S-F-B]	450	0	450

[A-S-RV-P-B]	418	0	418
[A-S-RV-P-C]	455	160	615
[A-S-RV-P-RV]	414	193	607

6

Since [A-S-RV-P-B-418] is the cheapest path in the frontier and passes the goal test, return it as the SOLUTION

When is A* Optimal?



S-G-5 passes the goal test
and returned as the SOLUTION

- A* returned S-G (actual cost 5) although S-A-G (actual cost 4) is the optimal solution
- Why?
 - Because A's heuristics is a bad estimate, it is much higher than the real cost of getting from A to G (which is 3)
- We need estimates (heuristics) to be **less** than actual costs!
estimated cost to goal $h(n) \leq$ actual cost to goal $h^*(n)$

Conditions for Optimality: Admissibility

- For A* search to be optimal, it needs an **admissible** heuristic, i.e.

$$0 \leq h(n) \leq h^*(n)$$

where $h^*(n)$ is the **true** cost of the **cheapest** solution from n to a nearest goal

- An admissible heuristic **never overestimates** the cost to reach the goal
- Since $g(n)$ is the actual cost to reach n along the current path, and $f(n) = g(n) + h(n)$, we can conclude that $f(n)$ never overestimates the true cost of a solution along the current path through n
- * Why is the straight-line distance, $h_{SLD}(n)$ an admissible heuristic for Romania?
- * What about the Manhattan distance for Pacman getting to a location in the maze?

A* Properties

- **Complete?** Yes, unless there are infinitely many nodes with $f < f(G)$
- **Optimal?** Yes – cannot expand f_{i+1} until f_i is finished
- **Time?** Exponential in [relative error of h x length of solution]
- **Space?** Keeps all nodes in memory. A* usually runs out of space long before it runs out of time \therefore not practical for many large-scale problems
- If C^* is the cost of the optimal solution path:

A* expands all nodes with $f(n) < C^*$

A* expands some nodes with $f(n) = C^*$

A* expands no nodes with $f(n) > C^*$

A* Applications

- Pathing problems
- Video games
- Robot motion planning
- Language analysis
- Machine translation
- Speech recognition

Heuristics for 8-puzzle

7	2	4
5		6
8	3	1

Start State


1	2	3
4	5	6
7	8	

Goal State


- Slide tiles horizontally or vertically into the empty space until the configuration matches the goal configuration
- Branching factor = number of possible moves at a particular state (maximum = 4, average = 3)
- Want to find shortest solutions by using A*:
 - ▶ h_1 = number of misplaced tiles
 - ▶ h_2 = total Manhattan distance (no. of squares from desired location of each tile or the sum of the distances of the tiles from their goal positions)
- Q: What are h_1 and h_2 for the Start State above?

Creating Admissible Heuristics

- Most of the work in solving hard search problems optimally is coming up with admissible heuristics
- Often admissible heuristics are solutions to relaxed problems where new actions are available



Pacman relaxed version: No wall




Relaxed Problems



- In 8-puzzle, what if the rules were relaxed so that a tile could move **anywhere** instead of just to an adjacent square?

$h_1(n)$ (number of misplaced tiles) gives the shortest solution

- If the rules were relaxed so that a tile could move to **any adjacent square**, then $h_2(n)$ (Manhattan distance) gives the shortest solution
- A relaxed problem is one with fewer restrictions on the actions (e.g. *simplified* version of the puzzle)
- $h(n)$ can be generated using machine learning techniques




A* Graph Search



Q: What went wrong here?

A* Graph Search



Q: What went wrong here?

A* graph search is not optimal because it doesn't expand $C(2+1)$ as it is already in the explored list via $C(3+1)$, so it fails to return the optimal solution S-A-C-G

Summary

- Informed search strategies may have access to **heuristic functions** $h(n)$ that estimate the cost of the shortest paths
- Good heuristics can dramatically reduce search cost
- **Greedy** best-first search expands the lowest h
 - ▶ incomplete and not always optimal
- **A*** expands lowest $g+h$ (UCS and Greedy)
 - ▶ complete and optimal (and optimally efficient)
- Admissible heuristics never overestimate the cost to the goal
- Heuristic design is key: can be constructed by relaxing the problem (e.g. 8-puzzle)

References

- Russel and Norvig, Chapter 3.5–3.6
- Greedy search quick demo [[Link](#)]
- A* search demo [[Link](#)]
- A* search explanation by John Levine [[Link](#)]
- Code in C++ for BFS, DFS, UCS, Greedy and A*, Y.Li [[Link](#)]