

Local Search:

4. Path does not Matter

Outline


- Recap – informed search & heuristics
- Local search/Iterative Improvement algorithms/Optimisers
 - ▶ hill-climbing (gradient ascent/descent)
 - ▶ simulated annealing
 - ▶ local beams
 - ▶ genetic algorithms

Recap – Informed Search

- Informed search strategies may have access to heuristic functions $h(n)$ that estimate the cost of the shortest paths
- Greedy best-first search expands the lowest h
 - ▶ incomplete and not always optimal
- A* expands lowest $g + h$
 - ▶ complete and optimal (and optimally efficient)
- **Admissible** heuristics never overestimate the cost to reach the goal

The Story So Far...


- The search algorithms we've seen so far explore search spaces **systematically** by keeping one or more paths in memory and by keeping track of which alternatives have been explored at each point along the path
- When a goal is found, the **path** to the goal also constitutes a **solution** to the problem
- But many problems do not care about the path, e.g. *8-queens*, just the **final configuration** and not which order the queens were added
- *Scheduling* problems (e.g. time allocation, resource allocation), just the final schedule matters, not how you arrived at that schedule
- In many *optimisation* problems, **path** is **irrelevant**, goal itself is the solution



Local Search



- Useable when solutions are states, not paths
- Start with a “complete” configuration (a single current state) and make modifications (**iterative improvements**) to improve its quality
- Given state space = set of complete configurations, find **optimal** configuration, (e.g. highest possible value or least cost) or find configuration satisfying **constraints**, (e.g. no two class at the same time in a timetable)
- Advantages of local search
 - they use very little memory (usually a constant amount)
 - can find reasonable solutions in large or infinite state spaces for which systematic algorithms are unsuitable
- Local search algorithms are very generic and have been applied successfully to many industrial problems

Example: n-Queens Problem



- **Iterative** improvement
 - start with one queen in each column (e.g. leftmost diagram)
 - move a queen along its column to reduce number of pairs of attacking queens (middle diagram)
 - repeat the previous step until h is minimum
- "We are what we repeatedly do." – Aristotle

State-space Landscape



- A landscape has both a **location** (defined by the state) and an **elevation** (defined by the value for the heuristic cost function or objective function)
- If elevation corresponds to a **cost**, the aim is to find the lowest valley – **global minimum**
- If elevation corresponds to an **objective function**, the aim is to find the **highest peak** – **global maximum**
- Sometimes have to go sideways or even backwards in order to make progress towards the solution

Hill-climbing Search (Gradient Descent/Ascent)

- Most basic local search algorithm – sometimes called greedy local search
- A loop that continually moves in the direction of increasing value, i.e. **uphill**
- Terminates when it reaches a **peak** where no neighbour has a higher value
- Does not maintain a search tree, data structure for current node only records the state and the objective function
- Does not look ahead beyond the intermediate **neighbours** of the current state

Hill-climbing Pseudocode

function HILL-CLIMBING(*problem*) **returns** a state that is a local maximum


current \leftarrow MAKE-NODE(*problem*.INITIAL-STATE)

loop do

neighbor \leftarrow a highest-valued successor of *current*


if *neighbor*.VALUE \leq *current*.VALUE **then return** *current*.STATE

current \leftarrow *neighbor*




- “Like climbing Everest in thick fog with amnesia”
- At each step the current node is replaced by the best neighbour; here it is the neighbour with the highest value
- If a **heuristic** cost estimate h is used, it would be the neighbour with the **lowest** h

8-queens using Hill-climbing



(a)





(b)

- Each state has 8 queens on the board, one per column
- Successors: all possible states generated by moving a single queen to square in the same column. Each state has $8 \times 7 = 56$ successor states (8 rows x 7 options per column)
- Heuristic cost function, h is the number of **pairs of queens** that are directly or indirectly **attacking** each other
- In (a), $h = 17$, the best moves from here are marked by the lowest numbers on the board by moving one queen within a column
- In 5 steps only, (b) can be reached
- In (b), $h = 1$, however, there are no further moves possible to decrease h , i.e. stuck at a **local minimum** (valley that is lower than is neighbouring states but higher than the global minimum)

Hill-climbing Problems

- Stuck at **local optimum** (maximum or minimum):
 - ▶ peak: higher than neighbours but lower than global maximum, no progress
 - ▶ ridge: oscillate from side to side
 - ▶ plateaux: a flat local maxima
- Potential solution: **random restart**
 - ▶ repeat the search from another starting point



Simulated Annealing

- Hill-climbing algorithm is **incomplete** – often fails to find a goal when one exists because they can get **stuck** at a local maximum or minimum
- Try combining hill-climbing with **random walk** – moving to a successor chosen at random from a set of successor states
- **Idea:** escape local maxima by allowing some “**bad**” moves (choose a worse neighbour) but **gradually decrease their size and frequency**
- Fun fact: the name annealing comes from the process used to harden metals and glass by heating them to a high temperature and then letting them cool slowly to reach low energy crystalline
- Start at a high intensity (e.g. high temperature) and then gradually reduce the intensity by lowering the temperature

Simulated Annealing Pseudocode* (Minimising)

function SIMULATED-ANNEALING(*problem, schedule*) **returns** a solution state

current \leftarrow *problem*.INITIAL

for *t* = 1 **to** ∞ **do**

T \leftarrow *schedule*(*t*)

if *T* = 0 **then return** *current*

next \leftarrow a randomly selected successor of *current*

$\Delta E \leftarrow$ VALUE(*current*) – VALUE(*next*)

if $\Delta E > 0$ **then** *current* \leftarrow *next*

else *current* \leftarrow *next* only with probability $e^{-\Delta E/T}$

$T \rightarrow \infty$: random walk (RW)
 $T \rightarrow 0$: like hill-climbing (HC)

$$P(\text{next}) = \frac{1}{1 + e^{-\Delta E/T}}$$

**This algorithm was taken from the 4th edition of the textbook*

- In the innermost loop, instead of picking the *best* move (HC), it picks a **random** move
- Assessment: if the move improves the situation, it is always accepted, otherwise it is accepted with some probability less than 1 (RW)
- The probability decreases exponentially with the “**badness**” of the move, given by ΔE , the probability also decreases as the temperature *T* goes down

SA Intuition

- When T is high (beginning), probability of accepting “bad” moves is high (RW)
- When T is low (end), probability of accepting “bad” moves is low, only moves uphill (HC)

How is T Decreased?

- Some possibilities
 1. $T_n = T_{n-1} - \text{constant}$
 2. $T_n = a T_{n-1}$ with $a < 1$
 3. $T_n = \text{constant} / (1 + n)$
 4. $T_n = \text{constant} / \ln(1 + n)$

When does SA Terminate?

- When the temperature becomes less than a given value
- When the number of iterations exceeds a given value
- When no improvement occurs after a given number of iterations

SA Example: Minimising Objective Function

- Stochastic hill-climbing, $f(current)=107$, $T=10$

$f(next)$	ΔE	$e^{-\Delta E/T}$	$P(next)$
80	27	0.067	0.94
100	7	0.497	0.67
107	0	1	0.5
120	-13	3.669	0.21

SA will accept this. *Why?*

$$P(next) = \frac{1}{1 + e^{-\Delta E/T}}$$

- The column $f(next)$ represents a set of candidate states that the current state can sample from at a time step
- Whether a randomly sampled candidate state will be accepted depends on ΔE and T

SA Example: Maximising Objective Function, f

- Stochastic hill-climbing, $f(current)=107$, $T=10$

$f(next)$	ΔE	$e^{\Delta E/T}$	$P(next)$
80	27	14.88	0.06
100	7	2.01	0.33
107	0	1	0.5
120	-13	0.27	0.78

$$\Delta E < 0$$



$$P(next) = \frac{1}{1 + e^{\Delta E/T}}$$

SA will accept this

Why?

- For maximising, ΔE checking and probability calculation are slightly different
- $\Delta E = f(current) - f(next)$ should be negative for "better" states

Local Maxima/Minima




Local Beam Search

- **Idea:** Keep k states instead of 1; choose top k of all their successors
- Begin with k randomly generated states. At each step all the successors of *all* k states are generated. If any one is a goal, algorithm halts. Otherwise select k best successors from the list and repeat
- **Problem:** Quite often, all k states end up on same local hill (like HC)
- **Stochastic** beam search: Instead of choosing best k from successor candidates, choose k successors at **random** (like SA), biased towards good ones → resembles natural selection; *successors* (offsprings) of a *state* (organism) populate the next generation according to its *value* (fitness)
- Special cases:
 - When $k = 1$, equivalent to HC
 - When $k = \infty$, equivalent to BFS

4-Queens Local Beam Search

$$k = 3$$

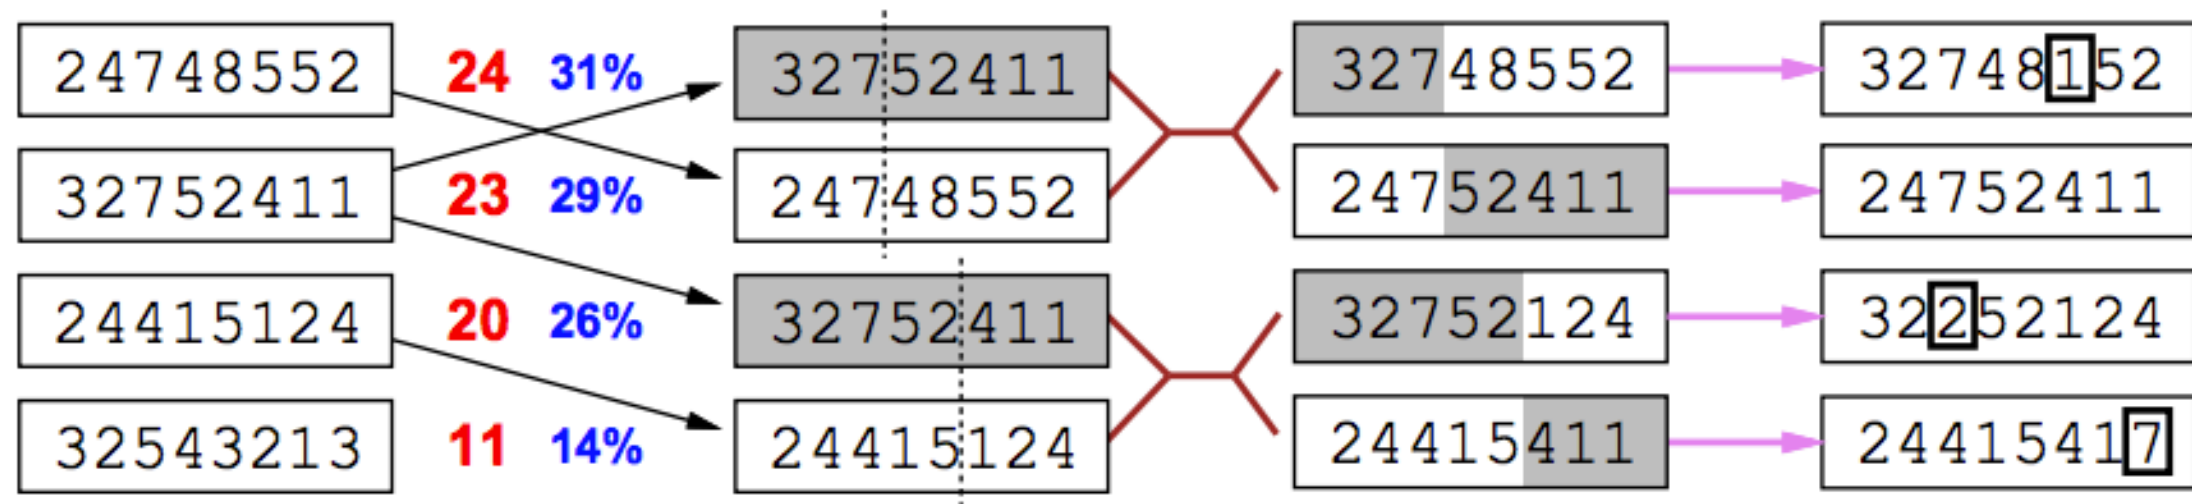
Until goal is found, choose "best" 3 states from all these successor states and repeat



Genetic Algorithms

- A variant of stochastic beam search in which successor states are generated by combining **two** parent states instead of one state
- Like beam searches, begin with a set of k randomly generated states, called the **population**
- Each state, or **individual** is represented as a string over a finite alphabet, most commonly 0s and 1s and is rated by the **fitness function** (objective function) – should return higher values for better states
- Based on this, pairs of individuals are **selected** and mated by choosing a **crossover** point from positions in the string
- Offsprings are created by crossing over the parent strings at the crossover point. Finally each location is subject to random **mutation** (e.g. in 8-queens, choosing a queen at random and moving it to a random square in its column)

GA Illustrated (8-queens)




- Population consists of individuals that are candidate solutions. Digits represent positions of queens on each column. Think of them as chromosomes.
- **Fitness** function: the number of *nonattacking* pairs of queens (28 for solution) = objective function
- Probability of being **selected** for reproduction (%) is directly proportional to the fitness score (= exploitation)
- **Pairs** – The individuals chosen are paired off at random; one of them has been chosen twice while one hasn't been chosen at all
- **Crossover** – when two parent states are quite different, produced states can be a long way from either parent, so crossover (like SA) frequently takes large steps in the state space early and smaller steps later on when most individuals are quite similar (= exploitation)
- **Mutation** – each location is subject to random mutation. In 8-queens, moving any queen to another position in the same column (= exploration)

Exploitation vs Exploration

- Exploitation – using already exist solutions and make refinement to it so its fitness will improve
 - ▶ Selection: use of solutions with high fitness to pass on to next generations
 - ▶ Crossover: main role is to provide mixing of the solutions and convergence in a subspace
- Exploration – the algorithm searches for new solutions in new regions
 - ▶ Mutation: the change of parts of one solution randomly, which increases the diversity of the population and provides a mechanism for escaping from a local optimum. Leads to a solution outside the subspace

GA: Pseudocode



* Question

What would a mutation in the 6th digit (to the string 32748**1**52) result in?

function GENETIC-ALGORITHM(*population*, *fitness*) **returns** an individual

repeat

weights ← WEIGHTED-BY(*population*, *fitness*)

population2 ← empty list

for *i* = 1 **to** SIZE(*population*) **do**

parent1, *parent2* ← WEIGHTED-RANDOM-CHOICES(*population*, *weights*, 2)

child ← REPRODUCE(*parent1*, *parent2*)

if (small random probability) **then** *child* ← MUTATE(*child*)

add *child* to *population2*

population ← *population2*

until some individual is fit enough, or enough time has elapsed

return the best individual in *population*, according to *fitness*

function REPRODUCE(*parent1*, *parent2*) **returns** an individual

n ← LENGTH(*parent1*)

c ← random number from 1 to *n*

return APPEND(SUBSTRING(*parent1*, 1, *c*), SUBSTRING(*parent2*, *c* + 1, *n*))

Only produces one offspring, not two, but *k* overall

Termination:

- $f(n)$ reached predefined value
- Cutoff time/generations
- No improvement in $f(n)$

Perform crossover between a pair and return one string/child

A genetic algorithm. Within the function, *population* is an ordered list of individuals, *weights* is a list of corresponding fitness values for each individual, and *fitness* is a function to compute these values.

GA: Worked Example

- Problem: maximise the function $f(x) = x^2$ where x is between 1 and 31
- Variables codes as strings. Numbers between 1 and 31 can be coded as 5-digit binary string, 00001_2 (1_{10}), 01010_2 (10_{10}), 11111_2 (31_{10})
- Start with population size 4, randomly
- Fitness/Objective function is x^2

Population: Iteration 0, Generation 0

Initial Population	x	f(n) = x ²
01101	13	169
11000	24	576
01000	8	64
10011	19	361

- Sum of $f(n)$, $\Sigma f = 169 + 576 + 64 + 361 = 1170$
- Average of $f(x) = 293$
- $\max(f(x)) = 576$

Selection

$f(x) = x^2$	$f(x)/\Sigma f$	$f(x)/$	x
169	0.14	0.58	13
576	0.49	1.97	24
64	0.06	0.22	8
361	0.31	1.23	19

- Many selection algorithms, e.g. Roulette wheel
- The best strings get more copies, weak ones die off
- After selection, crossover takes place
- Apply one bit mutation

Crossover & Mutation

Selected	String	Crossover	Mutation
13	01101	01100	01101
24	11000	11001	11101
19	10011	10000	10100
24	11000	11011	11001

- * At which point was crossover for 1st pair and 2nd pair?
- * Which bit was mutated in each child?


Next Generation, G1

New Population	x	$f(x) = x^2$	$f(x)/\Sigma f$
01101	13	169	0.08
11101	29	841	0.41
10100	20	400	0.2
11001	25	625	0.31


- $\Sigma f = 169 + 841 + 400 + 625 = 2035$
- Average of $f(x) = 508$
- $\max(f(x)) = 841$
- Is this population better than initial population?

Initial Population:
 $\Sigma f = 1170$
 $\text{mean}(f(x)) = 293$
 $\max(f(x)) = 576$

GA – Discussion




- Work well for continuous and discrete problems
- Tend not to get stuck in local maxima/minima (thanks to mutation)
- Computationally expensive
- Easy to perform in parallel
- Objective (fitness) function may be hard
- Practical applications: Knapsack problem, feature selection, travelling salesman (TSP)






- Both are nature-inspired
- Both can be used to solve same hard optimisation problems
- SA is a **single** state method, GA is a **population** method
- GA is normally better than SA performance-wise but at higher computation cost (can be distributed)
- SA might outperform GA in problems where solution space is small



Summary



- Hill-climbing: greedy local search that continually moves **uphill**; incomplete as can get stuck in local maxima
- Simulated annealing (non-deterministic/stochastic/randomised)
 - explore successors wildly **randomly** (High Temp)
 - as time goes by, explore less wildly (Cool down)
 - until there is a time where things settle (Cold)
- Local beam: keep k states instead of 1
- Genetic algorithms: local beam search + generate successor(s) from **pairs** of states – "survival of the fittest"
- SA and GA have been used in many **optimisation** problems – VLSI layout problems, factory scheduling, airline scheduling and other large-scale optimisation tasks

References

- Russel and Norvig, Chapter 4
- Paper on Simulated Annealing: Metropolis *et al.* (1953). "Equation of State Calculations by Fast Computing Machines", Journal of Chemical Physics. 21(6)
- Genetic algorithms explanation (Georgia Tech) [[Link](#)]
- The Knapsack Problem & Genetic Algorithms (Computerphile) [[Video](#)]
- GA demo – Learning to walk [[Link](#)]
- GA demo – Helping snakes find their way to food (G. Muric) [[Video](#)] [[Code](#)]