

Game Playing

Alpha–Beta Pruning

Minimax: Properties


- Depth-first traversal (branching factor b , depth m)
- **Complete?** Yes if tree is finite
- **Optimal?** Yes against an optimal opponent
- **Time complexity?** $O(b^m)$
- **Space complexity?** $O(bm)$
- Time cost is not practical for real games (Chess: $m \approx 100$, $b \approx 35$)

Minimax Problems


- The number of game states it has to examine is exponential in the depth of the tree
- Can we ignore some nodes?
- Can eliminate large parts of the tree using pruning
 - eliminating possibilities from consideration without having to examine them
 - allows us to ignore portions of the search tree that make no difference to the final choice

Minimax: Are There Shortcuts?

- **Minimax problem:** Number of game states it has to examine is exponential in the depth of the tree
- We can cut it in half by **pruning**, i.e. ignore portions of the search tree that makes **no difference** to the final choice (reduce the number of evaluations and branching)
- Calculation of optimal decision by considering what we **already know** at each point in the process could lead to minimax decision without evaluating some nodes
- Remove the nodes that don't have to be evaluated – removing redundancy



Node Evaluation



- $B \leq 3$ MIN – the biggest value that I can have is 3
- $A \geq 3$ MAX – the smallest value that I can have is 3

8. Alpha-Beta Pruning

- Main idea: If Player has a choice to move to node n for consideration, and if there is already a **better choice** of value for Player from previously processed nodes, then n can be **ignored**
- Alpha (α) is concerning what is the **minimum**, \geq that I can take (worst-case scenario), i.e. the first player who is trying to maximise the score
- Beta (β) is concerning what is the **maximum**, \leq that I have to give (worst-case scenario), i.e. the second player who is trying to minimise the score
- Pruning (termination of the recursive call) happens when the value of the current node is **worse than** the current alpha (Max) or beta (Min)

Alpha-Beta: Pseudocode

function ALPHA-BETA-SEARCH(*state*) **returns** an action
 $v \leftarrow \text{MAX-VALUE}(\text{state}, -\infty, +\infty)$
 return the *action* in ACTIONS(*state*) with value v

function MAX-VALUE(*state*, α , β) **returns** a utility value
 if TERMINAL-TEST(*state*) **then return** UTILITY(*state*)
 $v \leftarrow -\infty$
 for each a **in** ACTIONS(*state*) **do**
 $v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$
 if $v \geq \beta$ **then return** v
 $\alpha \leftarrow \text{MAX}(\alpha, v)$
 return v

function MIN-VALUE(*state*, α , β) **returns** a utility value
 if TERMINAL-TEST(*state*) **then return** UTILITY(*state*)
 $v \leftarrow +\infty$
 for each a **in** ACTIONS(*state*) **do**
 $v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$
 if $v \leq \alpha$ **then return** v
 $\beta \leftarrow \text{MIN}(\beta, v)$
 return v

- 2 main functions with 3 parameters –
 s (current state), α (best explored option for Max from root to s) and β (best explored option for Min from root to s) and an output value, v
- At the start node:
 - Smallest value that I can have,
 $\alpha = -\infty$
 - Biggest value that I can have,
 $\beta = +\infty$
- Keep track of alpha & beta globally and locally (lower and upper bounds)
- v is the value used to manipulate α and β and it is passed back from the function

$\alpha - \beta$ Pruning Walkthrough

function **Alpha-Beta-Search**(s) returns an action
 $s \leftarrow \text{root}$ # start from root
 $v \leftarrow \text{Minimax-ab}(s, -\infty, +\infty)$ # initialise alpha and beta
 return the action in Actions(s) with value v

function **Minimax-ab**(s, α, β) returns a value v
 if Terminal-test(s) then **return** Utility(s) # base case





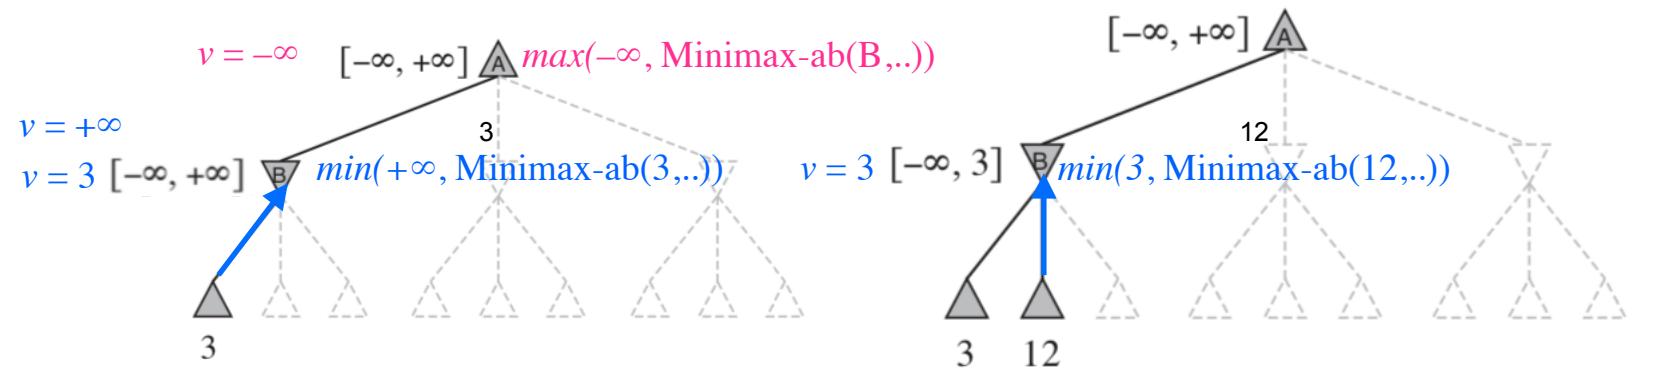
else if (Player == Max)

$v \leftarrow -\infty$ # reset v to $-\infty$ at every Max Player node
 for each child, c of s
 $v \leftarrow \max(v, \text{Minimax-ab}(c, \alpha, \beta))$
 if $v \geq \beta$ then **return** v # Pruning
 else $\alpha \leftarrow \max(\alpha, v)$

else if (Player == Min)

$v \leftarrow +\infty$ # reset v to $+\infty$ at every Min Player node
 for each child, c of s
 $v \leftarrow \min(v, \text{Minimax-ab}(c, \alpha, \beta))$
 if $v \leq \alpha$ then **return** v # Pruning
 else $\beta \leftarrow \min(\beta, v)$

return v



Alpha-beta Pruning Points

- α and β are **inherited** from the parent and they are manipulated locally at the current node for its own use
 - but they are not passed back up. Only v is passed back up (via return)
- At every node, v is **reset** to $-\infty$ if it is the node is a Max player or $+\infty$ if it is a Min Player
- At **Max**, **pruning** happens when v is **bigger than β** – Min Player above will ignore this path because it is worse (bigger) than the best that they have so far, β
- At **Min**, **pruning** happens when v is **smaller than α** – Max Player above will ignore this path because it is worse (smaller) than the best they have so far, α

Alpha–Beta: Properties

- Pruning does not affect final result
- Good move ordering improves effectiveness of pruning
- With perfect ordering, time complexity = $O(b^{\frac{m}{2}})$
 - ➔ doubles solvable depth
 - ➔ can easily reach depth 8 and play good chess

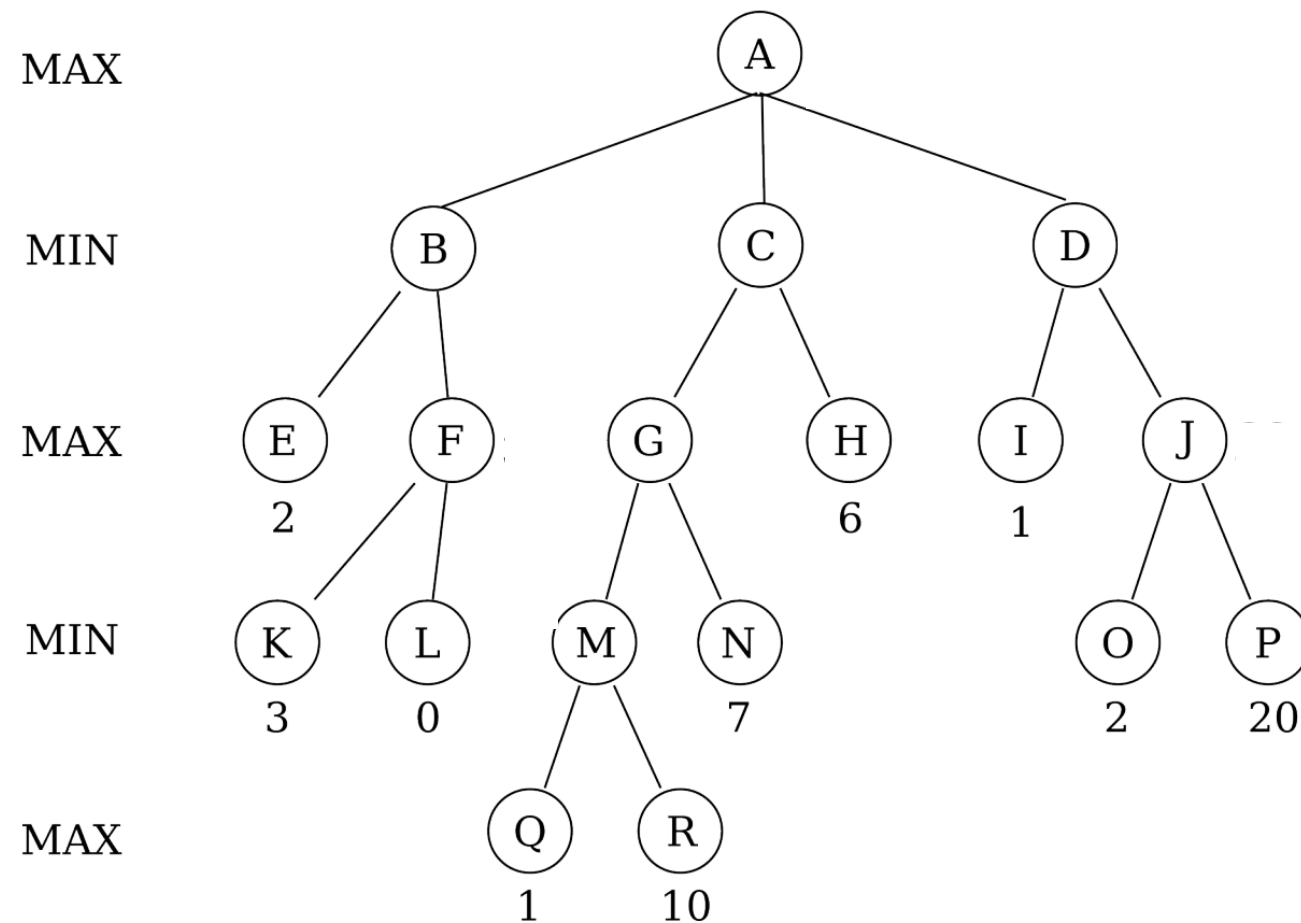
Summary

- Game playing as adversarial search
 - zero-sum games
 - utility values
- Search in games with perfect information:
 - Minimax
 - Alpha-beta pruning
 - ➔ Alpha-Beta has been used by popular programs like Deep Blue to efficiently play against Chess Grandmasters

References

- Russel and Norvig, Chapter 5, until 5.3
- J. Schrum, Alpha-beta pruning intuition [[Video](#)]
- S. Kambhapati, Alpha-beta intuition [[Video](#)]
- Historical reading:
 - Computer considers possible lines of play (Babbage, 1846)
 - Algorithm for perfect play (Zermelo, 1912; Von Neumann 1944)
 - Finite horizon, approximate evaluation (Zuse, 1945; Wiener, 1948; Shannon 1950)
 - First chess program (Turing, 1950)
 - Machine learning to improve evaluation accuracy (Samuel, 1952–57)
 - Pruning to allow deeper search (McCarthy, 1956)

Exercise: $\alpha - \beta$ Pruning



- Order the evaluations by nodes and α , β and v values
- First step:
 1. $A : \alpha = -\infty, \beta = +\infty, v = -\infty$

function **Alpha-Beta-Search**(s) returns an action
 $s \leftarrow \text{root}$ # start from root
 $v \leftarrow \text{Minimax-ab}(s, -\infty, +\infty)$ # initialise alpha and beta
 return the action in $\text{Actions}(s)$ with value v

function **Minimax-ab**(s, α, β) returns a value v
 if $\text{Terminal-test}(s)$ then **return** $\text{Utility}(s)$ # base case

else if (Player == Max)

$v \leftarrow -\infty$ # reset v to $-\infty$ at every Max Player node

for each child, c of s

$v \leftarrow \max(v, \text{Minimax-ab}(c, \alpha, \beta))$

if $v \geq \beta$ then **return** v # Pruning

else $\alpha \leftarrow \max(\alpha, v)$

else if (Player == Min)

$v \leftarrow +\infty$ # reset v to $+\infty$ at every Min Player node

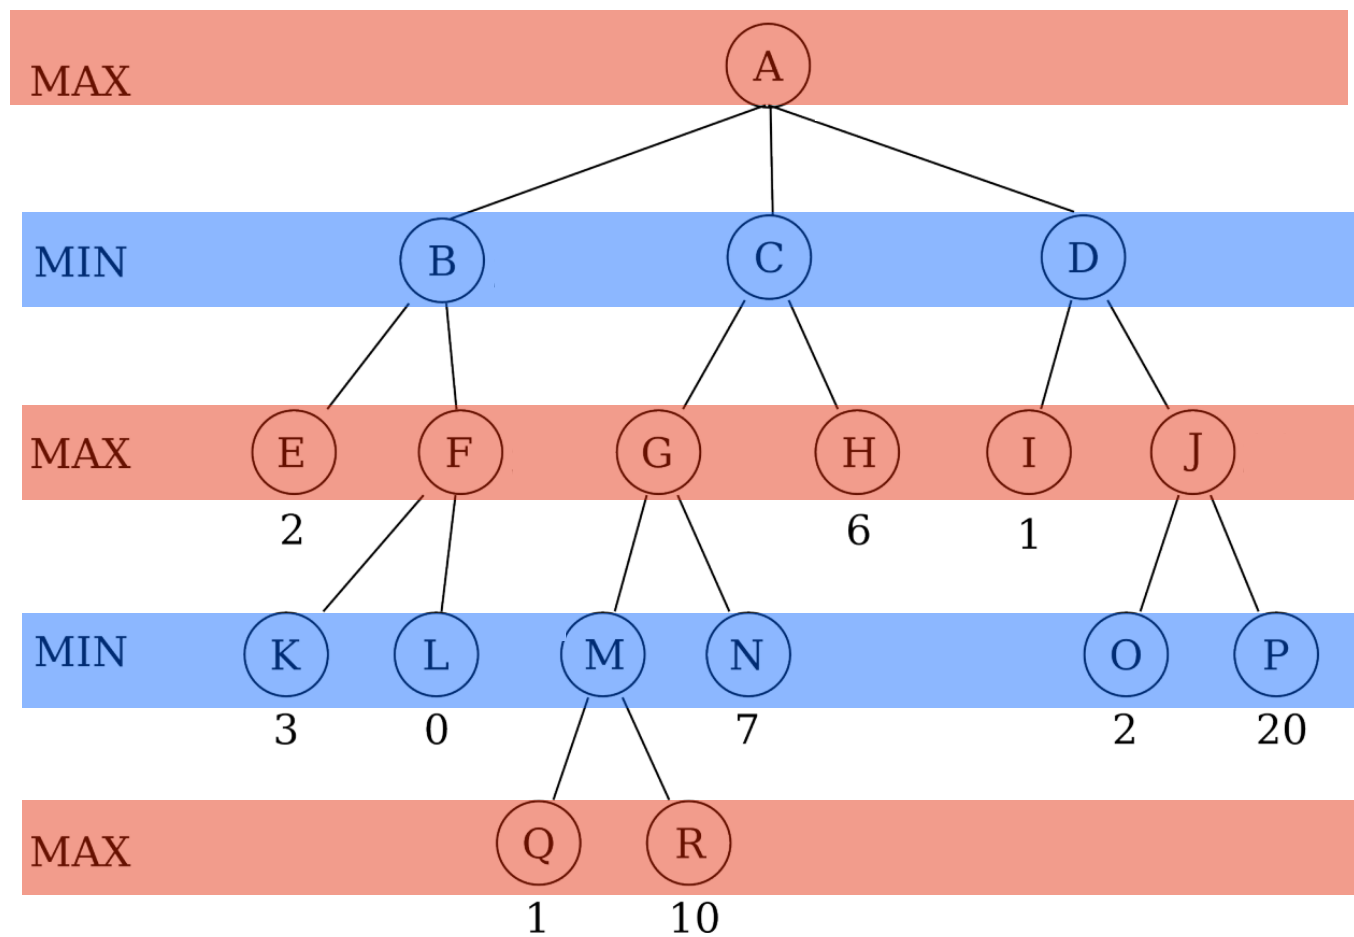
for each child, c of s

$v \leftarrow \min(v, \text{Minimax-ab}(c, \alpha, \beta))$

if $v \leq \alpha$ then **return** v # Pruning

else $\beta \leftarrow \min(\beta, v)$

return v



function **Alpha-Beta-Search**(s) returns an action
 $s \leftarrow \text{root}$ # start from root
 $v \leftarrow \text{Minimax-ab}(s, -\infty, +\infty)$ # initialise alpha and beta
 return the action in $\text{Actions}(s)$ with value v

function **Minimax-ab**(s, α, β) returns a value v
 if $\text{Terminal-test}(s)$ then **return** $\text{Utility}(s)$ # base case

else if ($\text{Player} == \text{Max}$)

$v \leftarrow -\infty$ # reset v to $-\infty$ at every Max Player node

for each child, c of s

$v \leftarrow \max(v, \text{Minimax-ab}(c, \alpha, \beta))$

if $v \geq \beta$ then **return** v # Pruning

else $\alpha \leftarrow \max(\alpha, v)$

else if ($\text{Player} == \text{Min}$)

$v \leftarrow +\infty$ # reset v to $+\infty$ at every Min Player node

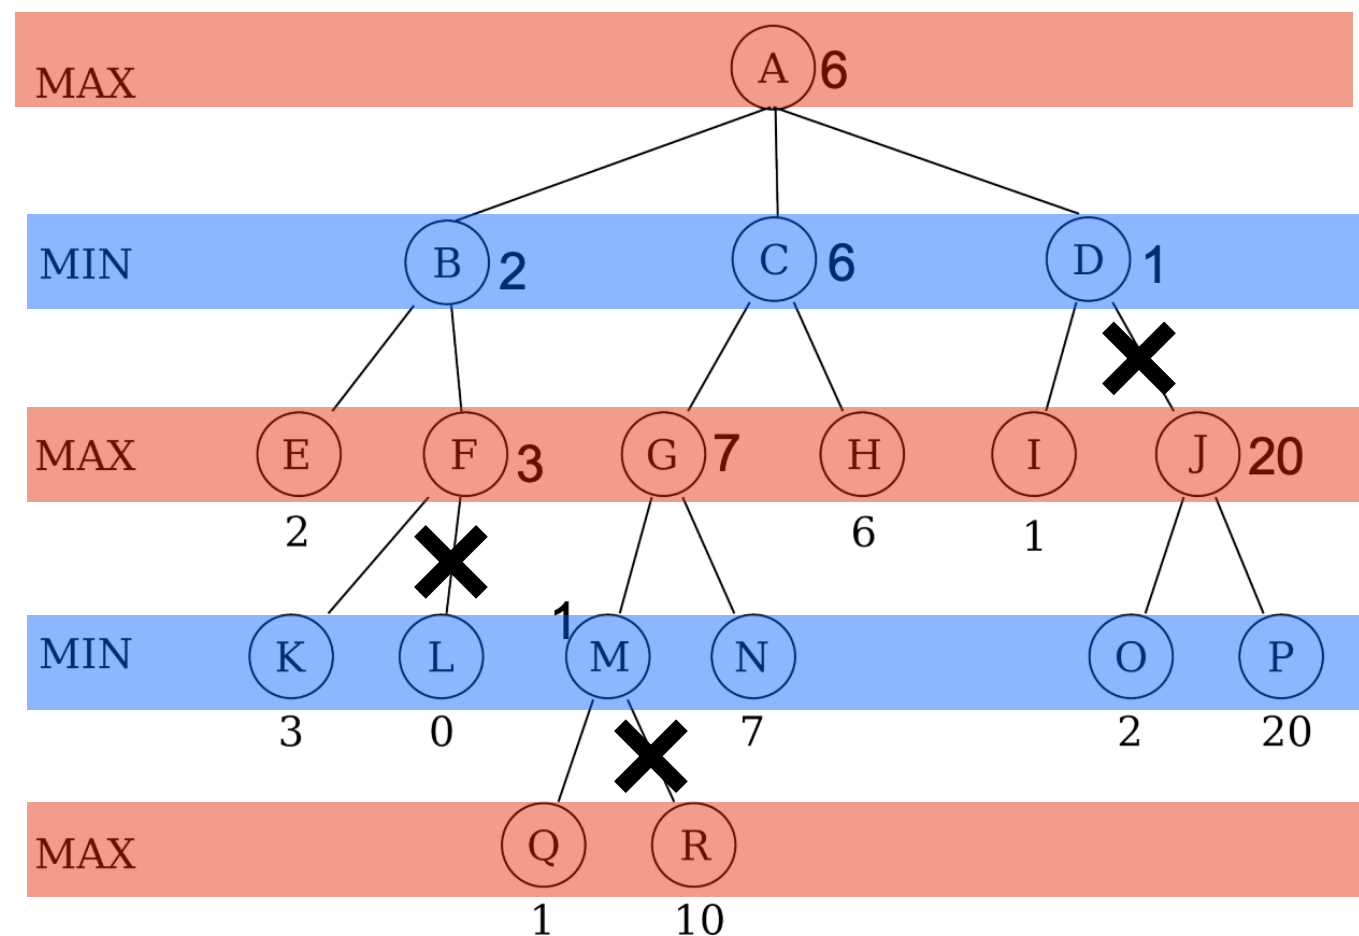
for each child, c of s

$v \leftarrow \min(v, \text{Minimax-ab}(c, \alpha, \beta))$

if $v \leq \alpha$ then **return** v # Pruning

else $\beta \leftarrow \min(\beta, v)$

return v



1. **A**: $\alpha=-\infty, \beta=+\infty, v=-\infty$

2. **B**: $\alpha=-\infty, \beta=+\infty, v=+\infty$

3. **E**: $\alpha=-\infty, \beta=+\infty, v=2$

4. **B**: $\alpha=-\infty, \beta=+\infty, v=+\infty$

5. **F**: $\alpha=-\infty, \beta=2, v=-\infty$

6. **K**: $\alpha=-\infty, \beta=2, v=3$

7. **F**: $\alpha=-\infty, \beta=2, v=-\infty$ (L ignored)

8. **B**: $\alpha=-\infty, \beta=2, v=2$

9. **A**: $\alpha=-\infty, \beta=+\infty, v=-\infty$

10. **C**: $\alpha=2, \beta=+\infty, v=+\infty$

11. **G**: $\alpha=2, \beta=+\infty, v=-\infty$

12. **M**: $\alpha=2, \beta=+\infty, v=+\infty$

13. **Q**: $\alpha=2, \beta=+\infty, v=1$

14. **M**: $\alpha=2, \beta=+\infty, v=+\infty$ (R ignored)

15. **G**: $\alpha=2, \beta=+\infty, v=-\infty$

16. **N**: $\alpha=2, \beta=+\infty, v=7$

17. **G**: $\alpha=2, \beta=+\infty, v=1$

18. **C**: $\alpha=2, \beta=+\infty, v=+\infty$

19. **H**: $\alpha=2, \beta=7, v=6$

20. **C**: $\alpha=2, \beta=7, v=7$

21. **A**: $\alpha=2, \beta=+\infty, v=2$

22. **D**: $\alpha=6, \beta=+\infty, v=+\infty$

23. **I**: $\alpha=6, \beta=+\infty, v=1$

24. **D**: $\alpha=6, \beta=+\infty, v=+\infty$ (J, O, P ignored)

25. **A**: $\alpha=6, \beta=+\infty, v=6$

Q: Which move should Max take? $A \rightarrow C$

Q: Which move should Min take after that? $C \rightarrow H$

Poll Results

How would you describe the course so far?

